



AI赋能的查询优化

南方科技大学

唐 博

tangb3@sustech.edu.cn





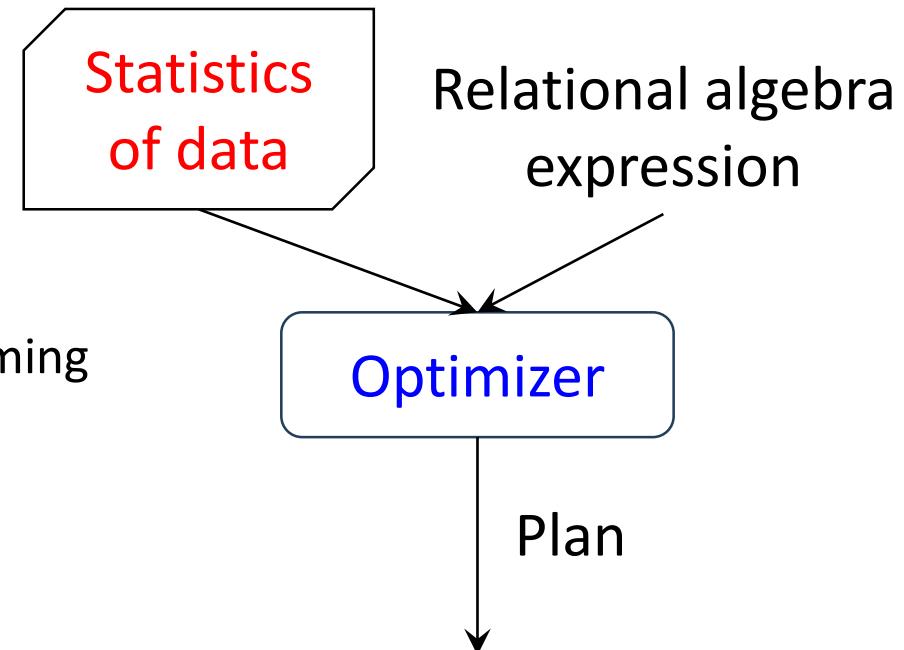
Recap: Why Query Optimizer?



❖ **Goal:** Find the fastest plan.

❖ **Techniques**

- Equivalence rules
- Join ordering by dynamic programming
- Cost and size estimation



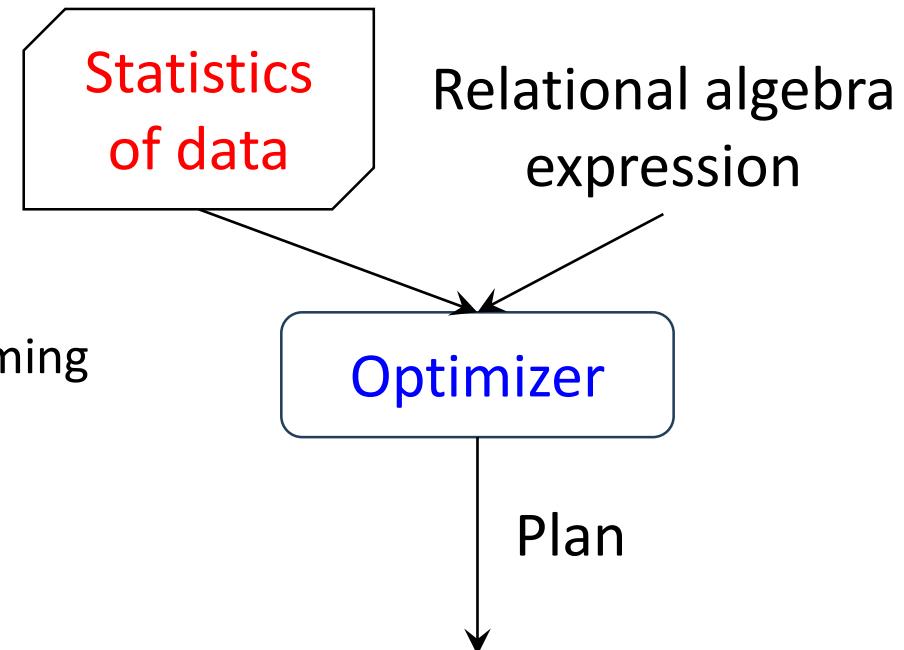
Recap: Why Query Optimizer?

❖ **Goal:** Find the fastest plan.

❖ **Techniques**

- Equivalence rules
- Join ordering by dynamic programming
- Cost and size estimation

Estimation always yields errors,
thus bad plan selection.





Recap: Estimate Selection Size



❖ **Problem:** Let $size$ be the estimated number of tuples satisfying the condition, estimate the size of $\sigma_{A=v}(r)$ and A is not a unique attribute.

- Estimated size = $n_r / V(r, A) = 8/4 = 2$
- Actual size = 1

❖ Estimate the size of $\sigma_{m \leq A \leq v}(r)$

- If no histogram,
 - Estimated size = $\frac{v-m}{\max(r,A)-\min(r,A)} = \frac{22-13}{26-1} = 0.36$
 - Actual size = 2

Relation r

ID	A
1	a
2	a
3	m
4	v
5	z
6	z
7	z
8	z

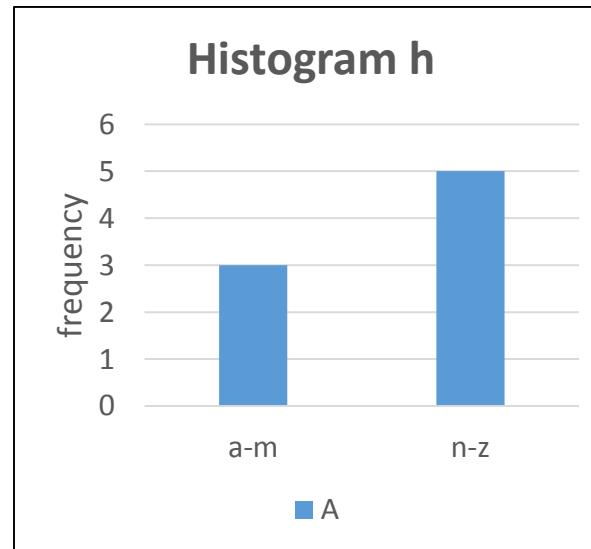
Recap: Estimate Selection Size

- ❖ Estimate the size of $\sigma_{m \leq A \leq v}(r)$

- If with histogram,

- Estimated size = $3*(1/13)+5*(9/13) = 3.69$

- Actual size = 2





Recap: Estimate Join Size



<i>Case of attributes</i>	<i>Size of r \bowtie s</i>	<i>Remarks</i>
...
$R \cap S = \{A\}$ is not a key for r nor s	$\frac{n_r * n_s}{\max\{V(r, A), V(s, A)\}}$	The lower estimate, assume that every tuple t in r (or s) produces tuples in $r \bowtie s$

❖ **Problem:** Estimate the size of $r \bowtie_{r.A=s.A} s$

- Estimated size = $\frac{4 * 4}{4} = 4$
- Actual size = 3

Relation r

ID	A
1	a
2	b
3	c
4	d

Relation s

ID	A
1	a
2	b
3	d
4	e



Recap: Estimate Join Size



❖ **Problem:** Generate join order for three relations r, s, t

- $r(A) \bowtie_A s(A) \bowtie_A t(A)$, Each relation has 1000 tuples.
- For simplicity, assume the cost of a join operation
= the # of tuples of input intermediate relation(s).

	$\{r, s\}$	$\{r, t\}$	$\{s, t\}$		$\{r, s, t\}$
Estimated Size	1000	2000	500	Estimated Size	10000
Actual Size	100	1000	500	Actual Size	1000
Estimated Cost	0	0	0	Estimated Cost	500
Estimated best plan	$r \bowtie s$	$r \bowtie t$	$s \bowtie t$	Estimated best plan	$r \bowtie (s \bowtie t)$
Actual best plan	$r \bowtie s$	$r \bowtie t$	$s \bowtie t$	Actual best plan	$(r \bowtie s) \bowtie t$



Recap: Summary

❖ Cause of estimate error

- Skew distribution of practical data v.s. uniform distribution assumption (for selection).
- Lack of correlation statistics between columns intra and inter-tables (for join).

❖ Impact of estimate error

- worse plan selection → more time, more cost, worse user experience.



How to Improve Estimation?



- ❖ Improve estimation is to represent data distribution better.
 - For statistics-based methods, better representation → more statistical data (could be more than data itself).
 - Fortunately, **AI does great in learning data distribution.**



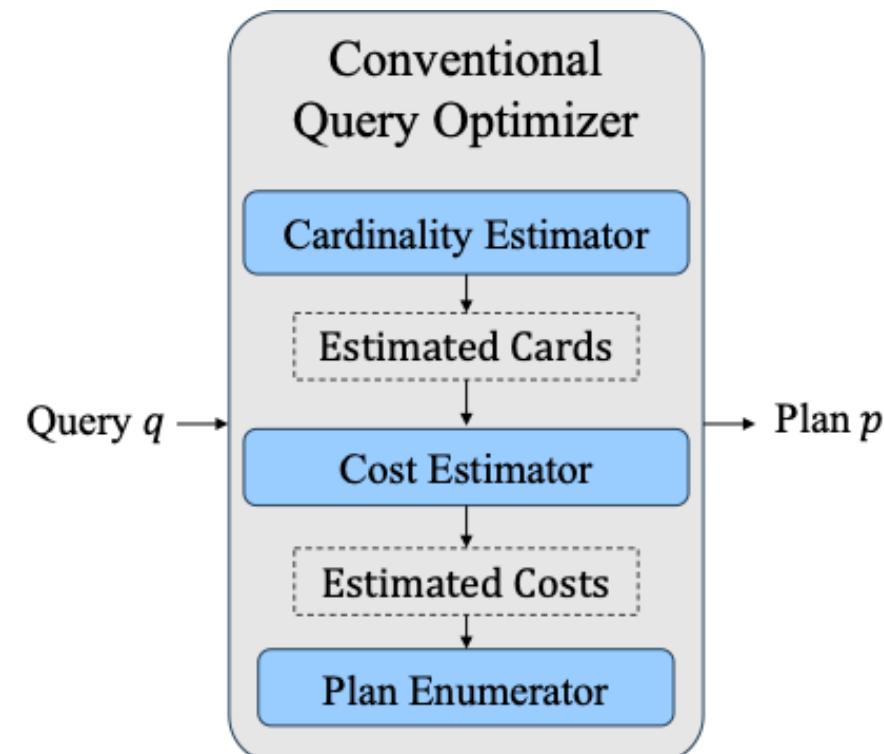
AI-based Query Optimization



- ❑ Learned Query Optimization
 - Learned Estimator
 - Learned Optimizer
 - Learned Optimizer Enhancer
- ❑ Other AI4DB Topics

❖ Components of Query Optimizer

- Cardinality Estimator: estimate operator (Selection/Join) size.
- Cost Model: estimate plan cost.
- Plan Enumerator: build plans via heuristic methods (e.g., Dynamic Programming).

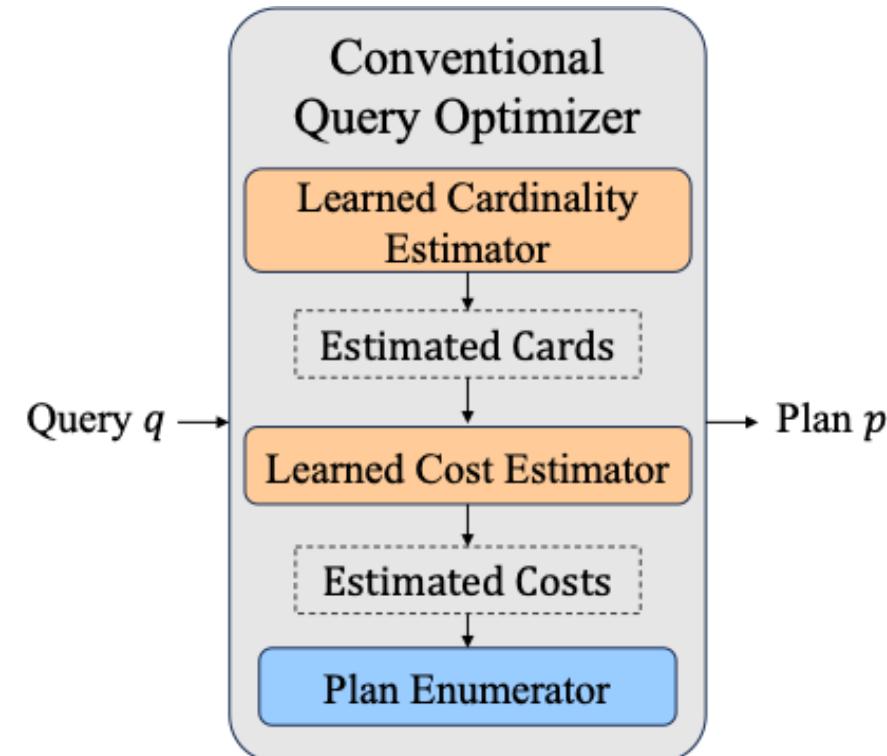


Simplified architecture of query optimizer

Learned Estimator

❖ Modules of Learned Estimator

- Learned Cardinality Estimator:
 - estimate operator (Selection/Join) size via a learning model.
- Learned Cost Model: estimate plan cost via a learning model.
- Plan Enumerator: no change.



Architecture of learned estimator



Learned Estimator



- Learned Cardinality Estimator
 - Query-driven: learn a mapping from query to cardinality from the historical workload.
 - Data-driven: learn the distribution directly from data.
 - Hybrid: combine both above methods.

- Learned Cost Model
 - Learn a mapping from plan to cost from the historical workload.



Learned Cardinality Estimator



❖ Query-driven (MSCN [1])

Step 1 - Query Featurization

```
SQL: SELECT COUNT(*) FROM title t, movie_companies mc WHERE t.id = mc.movie_id AND t.production_year > 2010 AND mc.company_id = 5;
```

↓ Transform to three sets of feature vector

Table set
 $\{[0 \underline{1} 0 \dots 0], [0 0 \underline{1} 0 \dots 1]\}$

table id
(one hot)

Join set
 $\{[0 \underline{0} 1 0]\}$

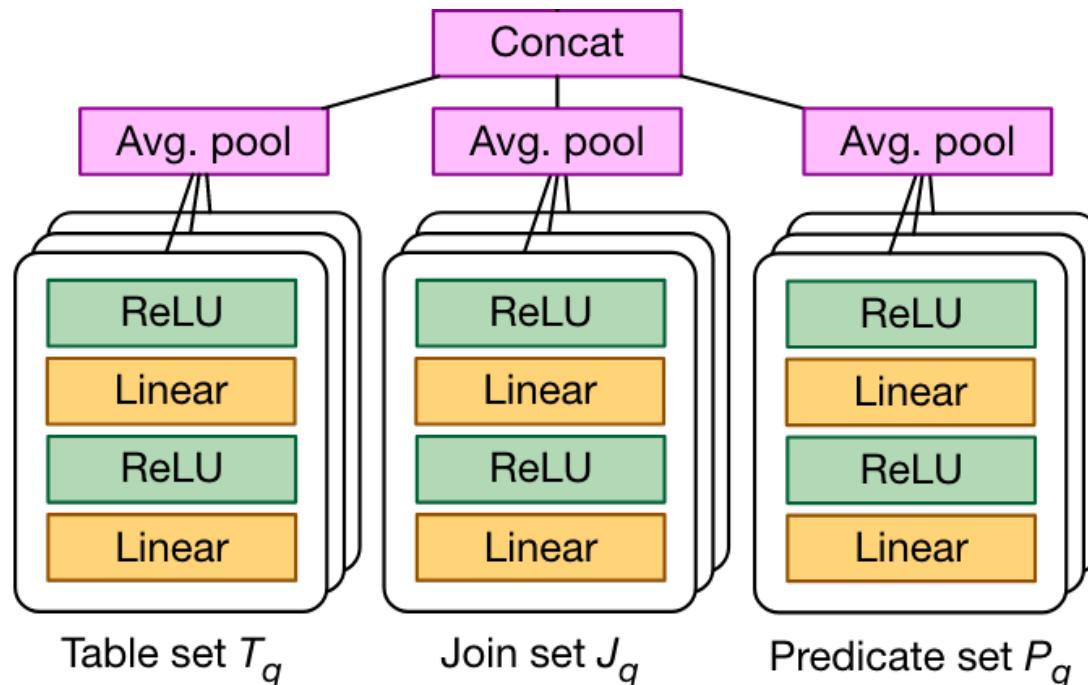
samples
join id
(one hot)

Predicate set
 $\{[\underline{1} 0 0 0 \underline{0} 1 0 0 \underline{0.72}], [\underline{0} 0 \underline{0} 1 0 0 \underline{1} 0 \underline{0.14}]\}$

column id
(one hot) value
(normalized) operator id
(one hot)

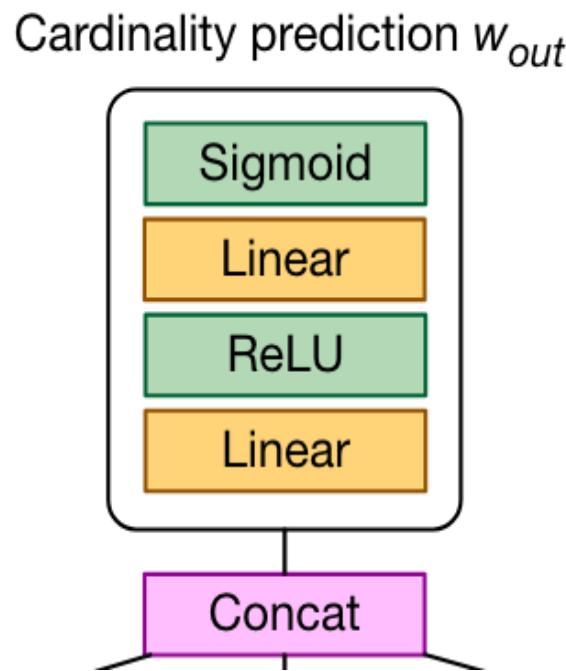
❖ Query-driven (MSCN [1])

Step 2 - Query Representation



- ❖ Query-driven (MSCN [1])

Step 3 - Cardinality Prediction





Learned Cardinality Estimator

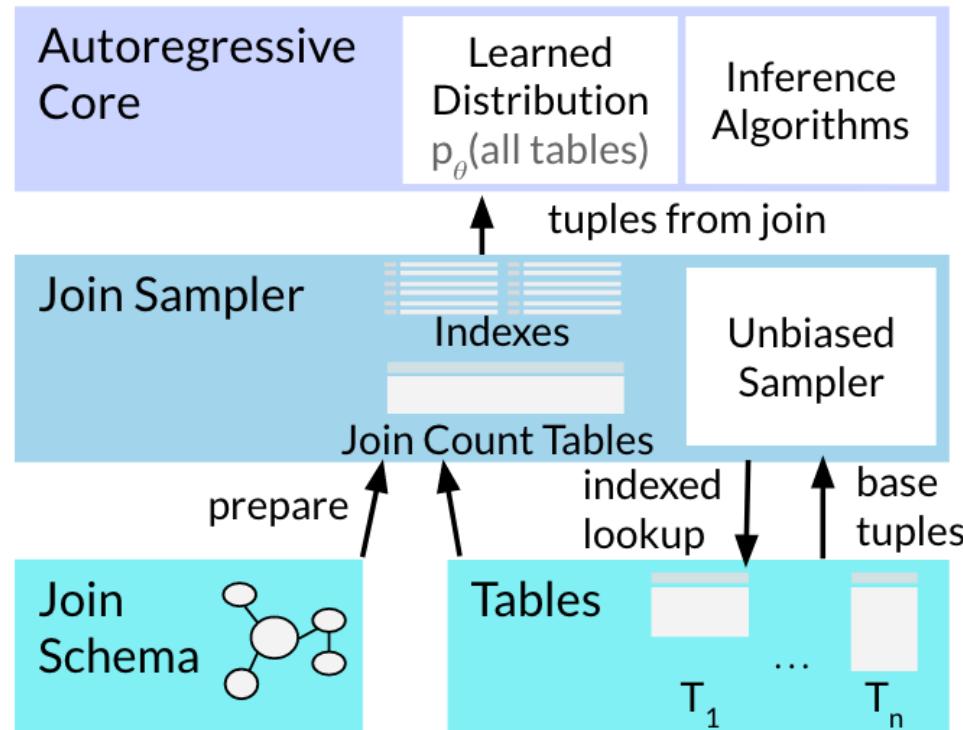


❖ Query-driven (MSCN [1])

Estimation error on the JOB-light workload

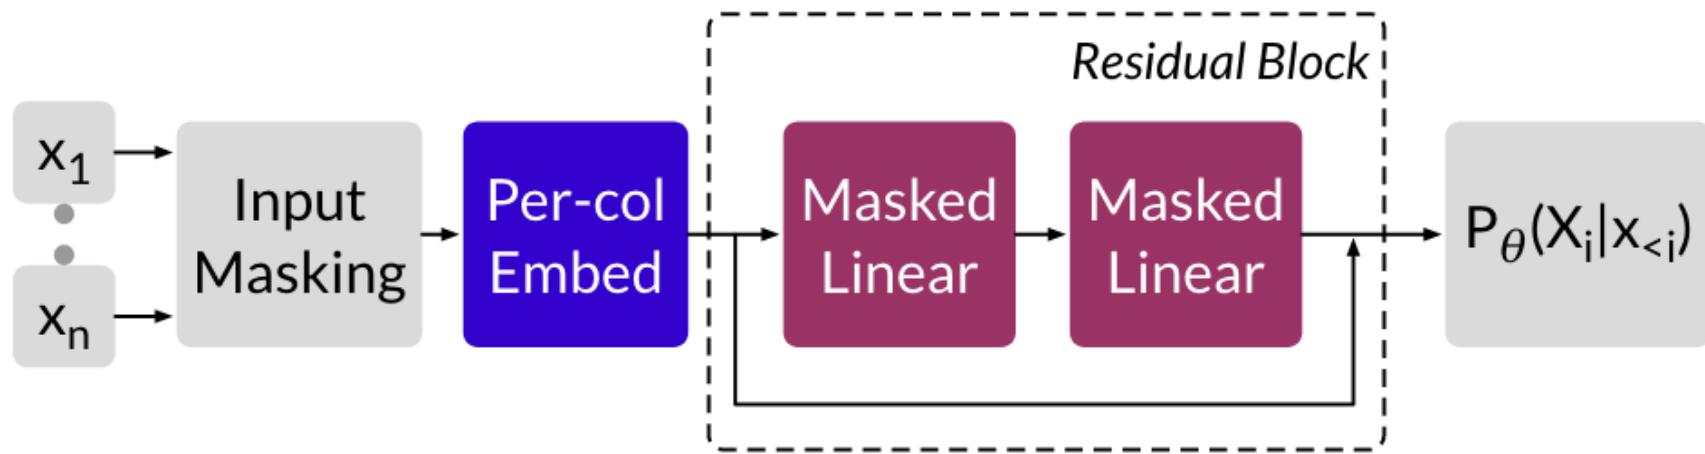
	median	90th	95th	99th	max	mean
PostgreSQL	7.93	164	1104	2912	3477	174
MSCN	3.82	78.4	362	927	1110	57.9

❖ Data-driven (NeuroCard [1])



Overview of NeuroCard

- ❖ Data-driven (NeuroCard [1])



Default architecture of the autoregressive model



Learned Cardinality Estimator



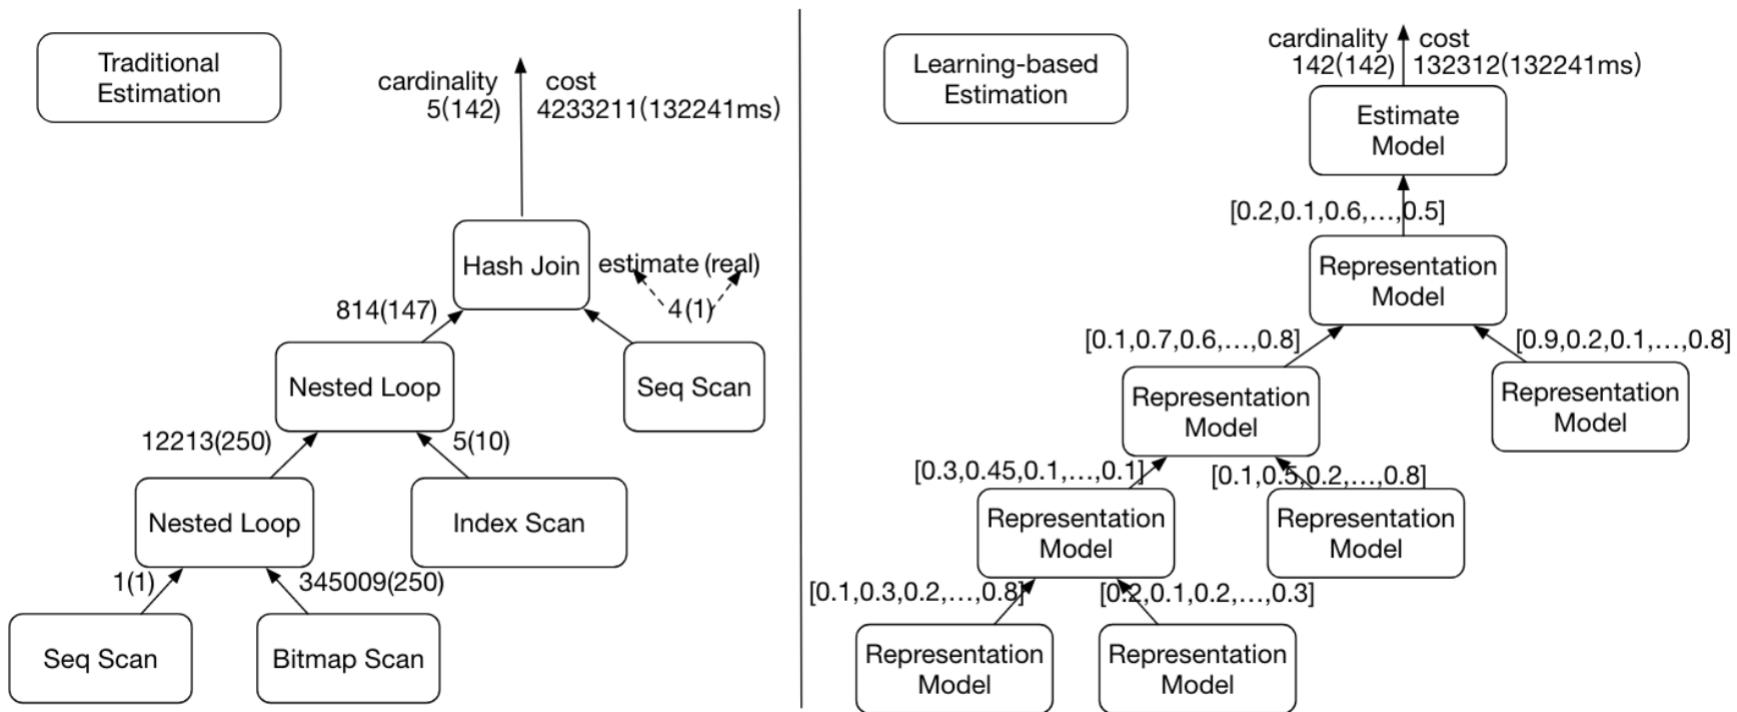
- ❖ Data-driven (NeuroCard [1])

Estimation error on the JOB-light workload

	median	95th	99th	max
PostgreSQL	7.97	797	$3 \cdot 10^3$	10^3
MSCN	3.01	136	$1 \cdot 10^3$	10^3
NeuroCard	1.57	5.91	8.48	8.51

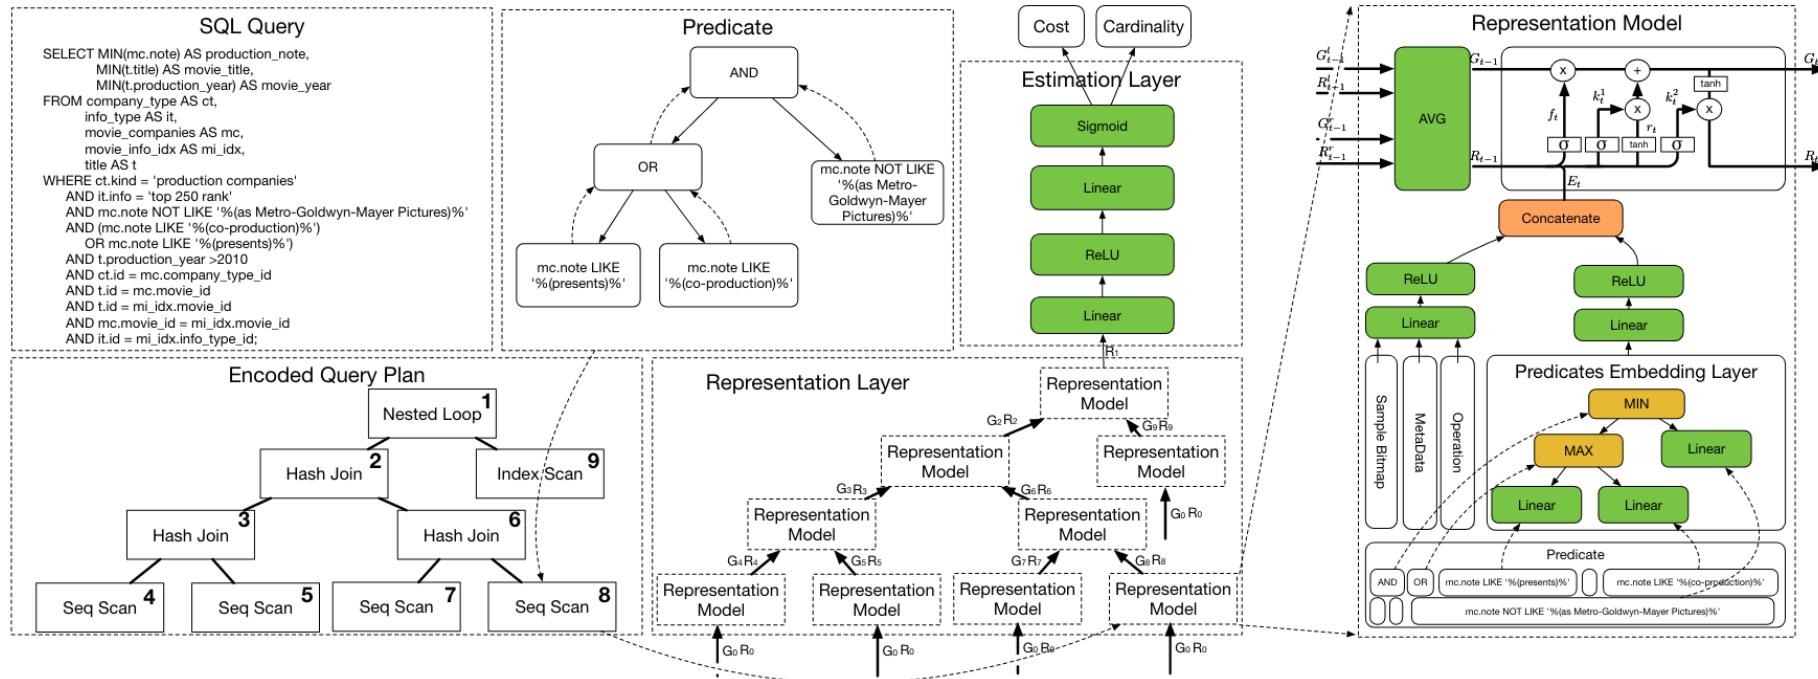
Learned Cost Model

- ❖ T-LSTM[1]: both cardinality and cost are inaccurate



Comparison of traditional cost estimation and learning-based cost estimation

❖ T-LSTM[1]: estimate both cardinality and cost



Two level tree model



Learned Estimator Analysis



- ❖ Learning historical workload or data for better estimation
- ❖ Pros:
 - Clear task definition and optimization procedure
- ❖ Cons:
 - Better estimations do not guarantee superior performance
 - Hard to integrate into DBMS
 - Significant overhead for multi-join queries
 - Insufficient robustness



Learned Optimizer



How to design a learned-based optimizer to improve
the overall query processing performance?



Learned Optimizer



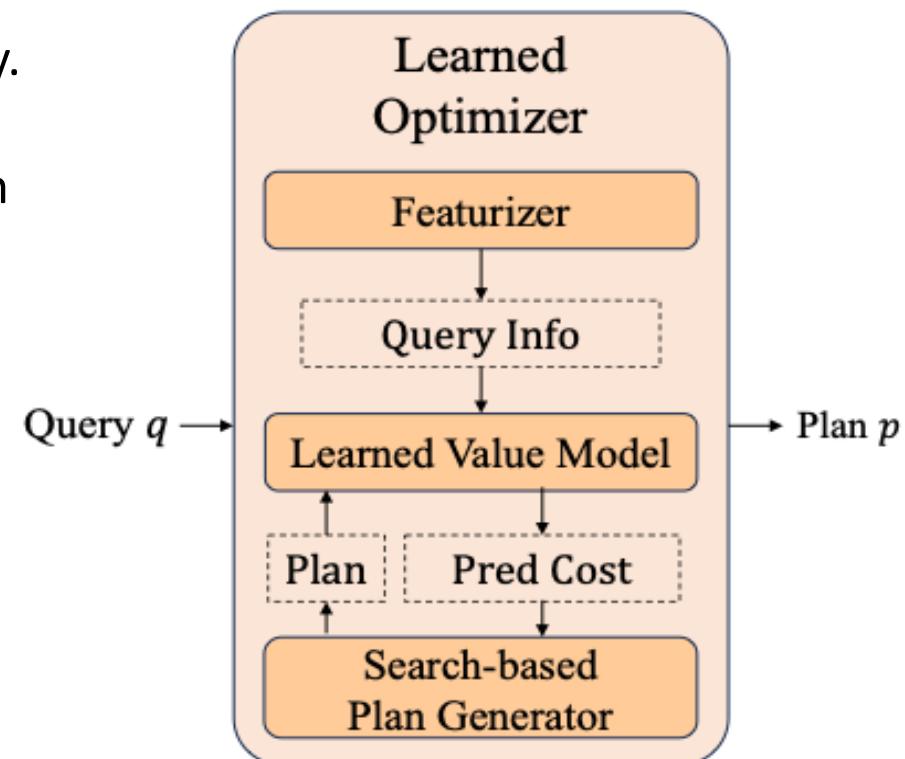
How to design a learned-based optimizer to improve
the overall query processing performance?

Design a reinforcement learning-based optimizer that
learns from its plan construction decision!

Learned Optimizer

❖ Components of Learned Optimizer

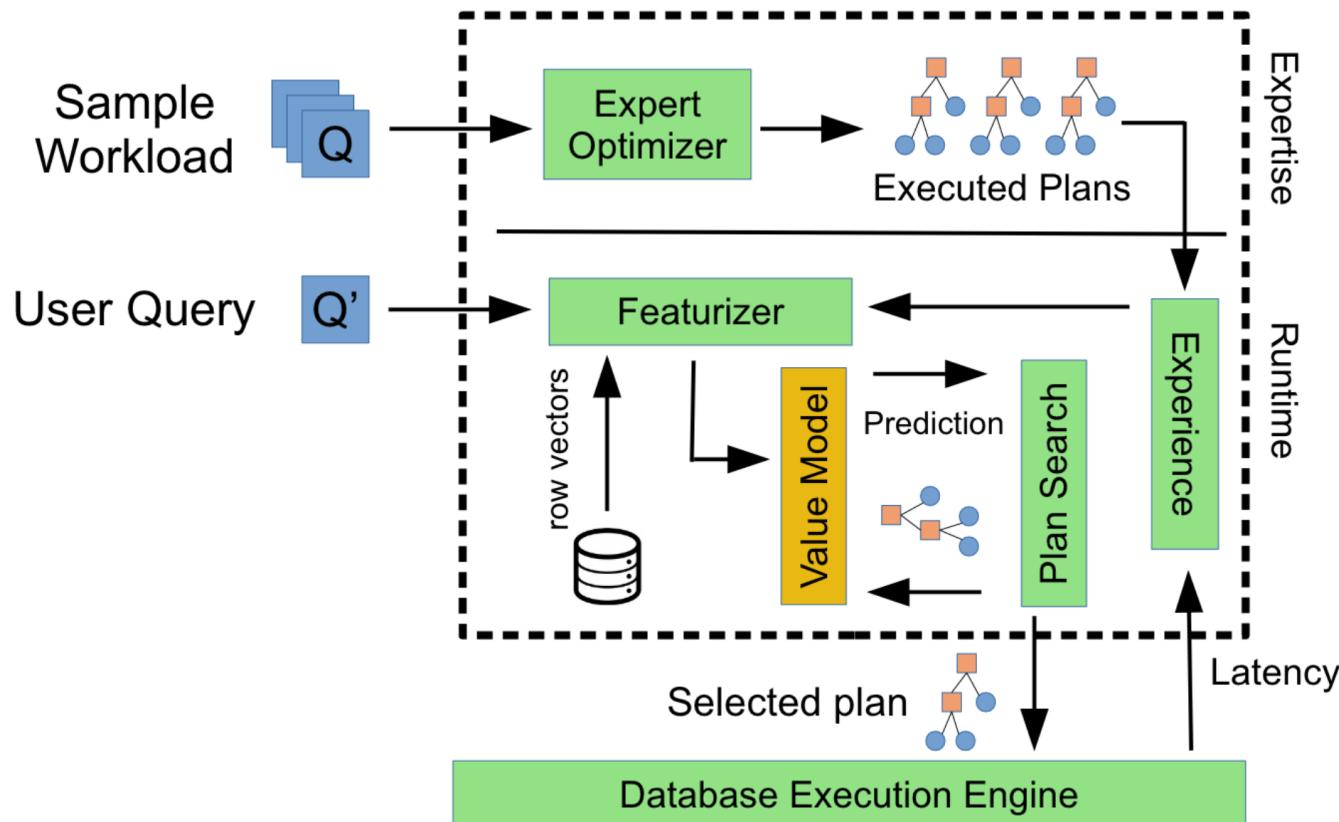
- **Featurizer**: extract features of query.
- **Learned Value Model**: evaluate plan construction decision.
- **Search-based Plan Generator**: efficient heuristic method to construct a plan (e.g., best-first search, beam search).



Architecture of learned optimizer

Learned Optimizer

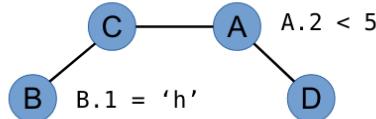
❖ Neo [1]



Architecture of Neo

❖ Neo [1] - Featurizer

SELECT * FROM A, B, C, D WHERE
 A.3=C.3 AND A.4=D.4 AND C.5=B.5
 AND A.2<5 AND B.1='h';



	A	B	C	D	E
A	0	0	1	1	0
B	0	0	1	0	0
C	1	1	0	0	0
D	1	0	0	0	0
E	0	0	0	0	0

Join Graph

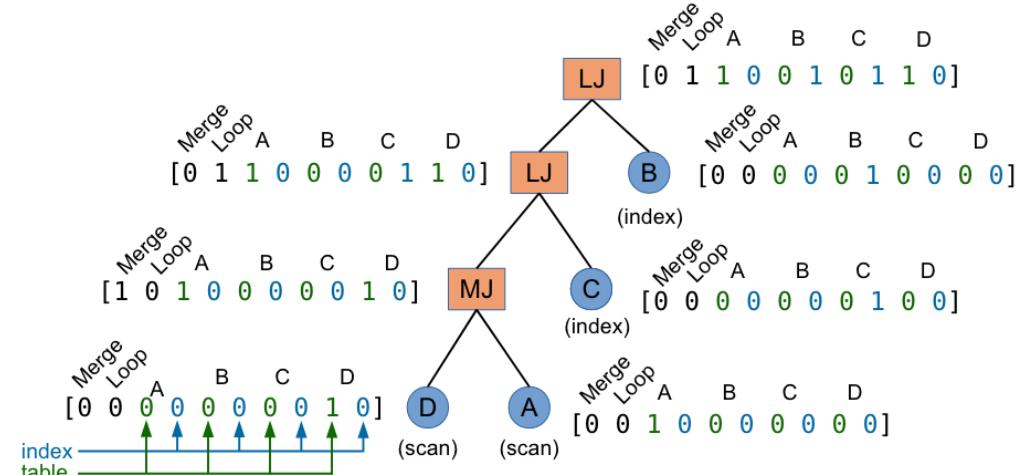
A.1	A.2	...	B.1	B.2	...	E.1	E.2
0	1	...	1	0	...	0	0

Column Predicates

[0 1 1 0 1 0 0 0 0 0 1 ... 1 0 ... 0 0]

Query-level Vector

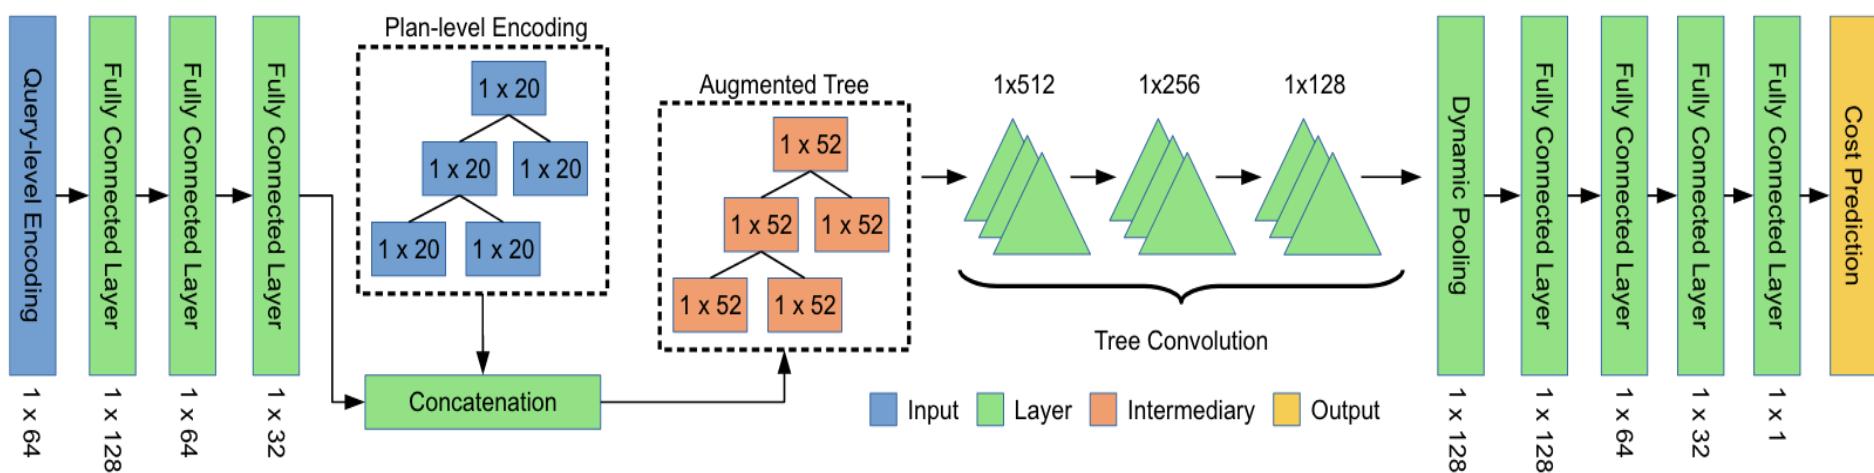
Query-level encoding



Plan-level encoding

Learned Optimizer

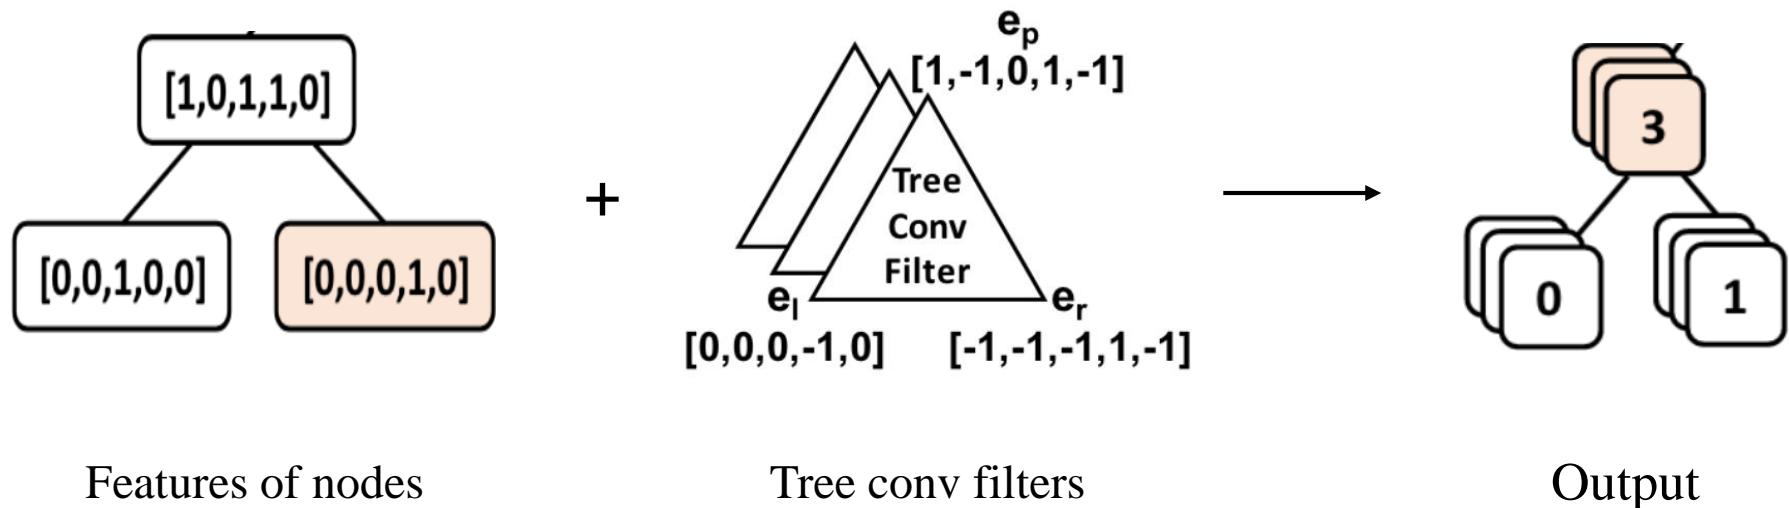
❖ Neo [1] - Value Network



Architecture of Neo's value network

Learned Optimizer

- ❖ Neo [1] - Value Network - Tree Convolution Example





Learned Optimizer Analysis



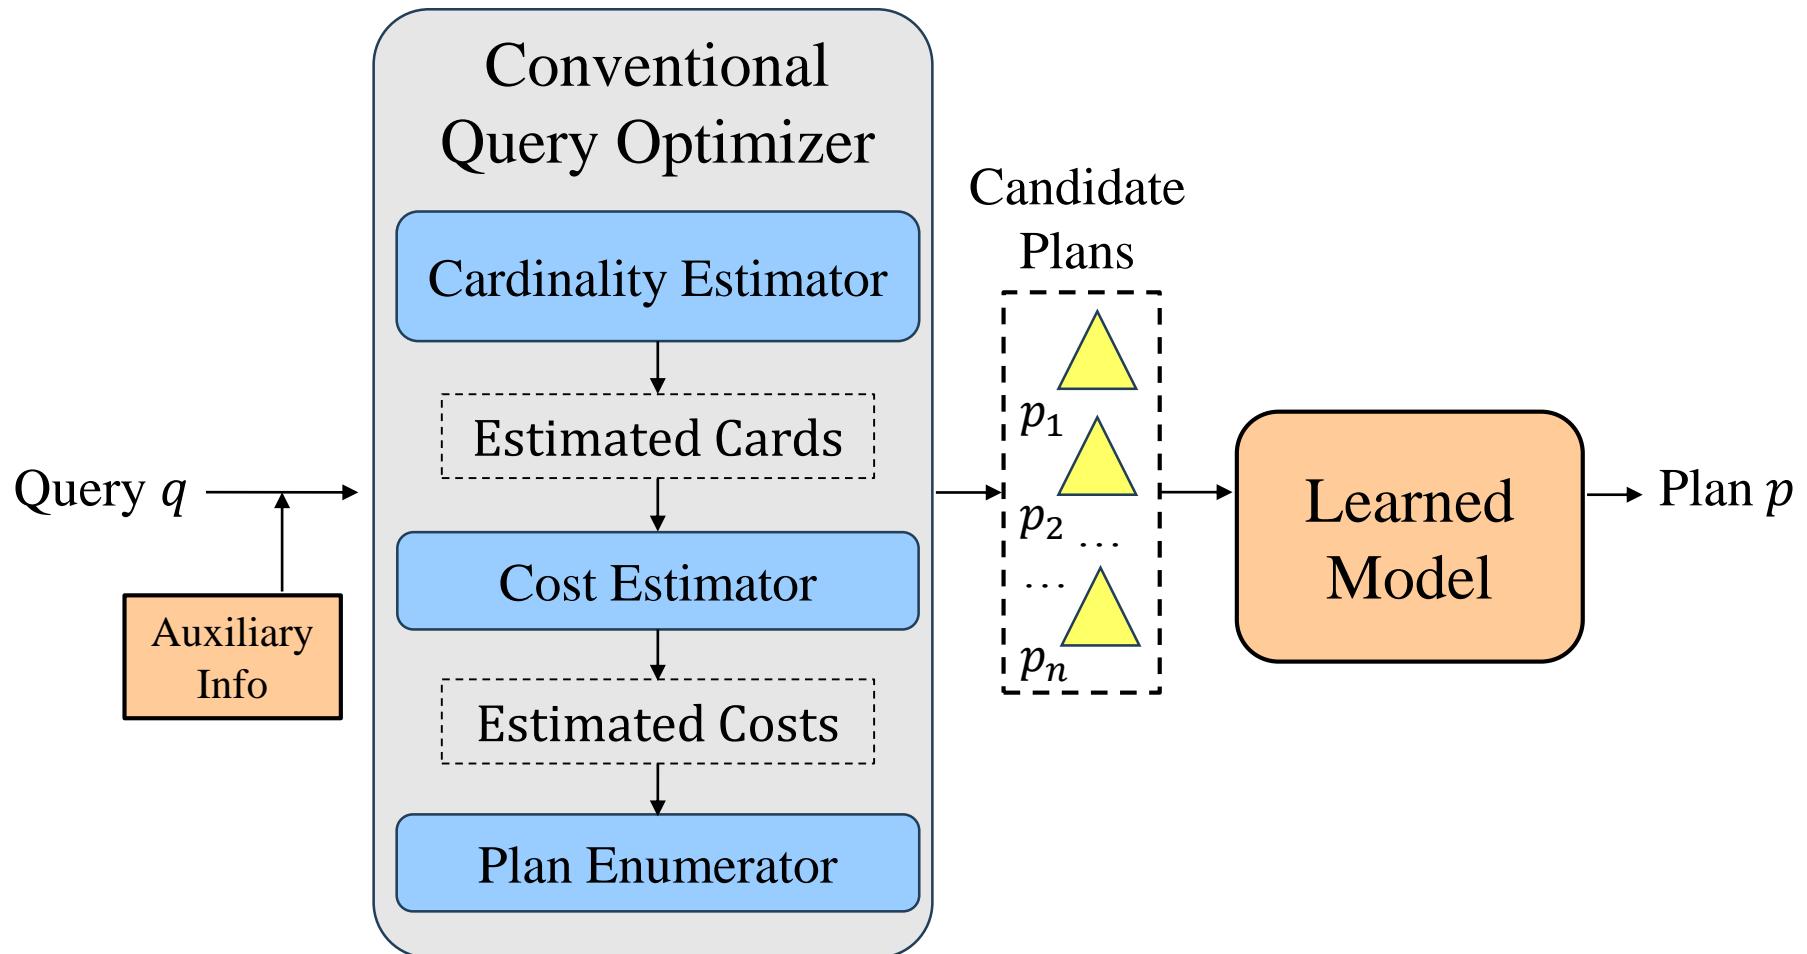
- ❖ Learn to construct plans via reinforcement learning
- ❖ Pros:
 - High potential for performance improvement
 - Available for database without a query optimizer
- ❖ Cons:
 - Substantial training overhead
 - Risk of worse plan generation (Insufficient robustness)



Learned Optimizer Enhancer

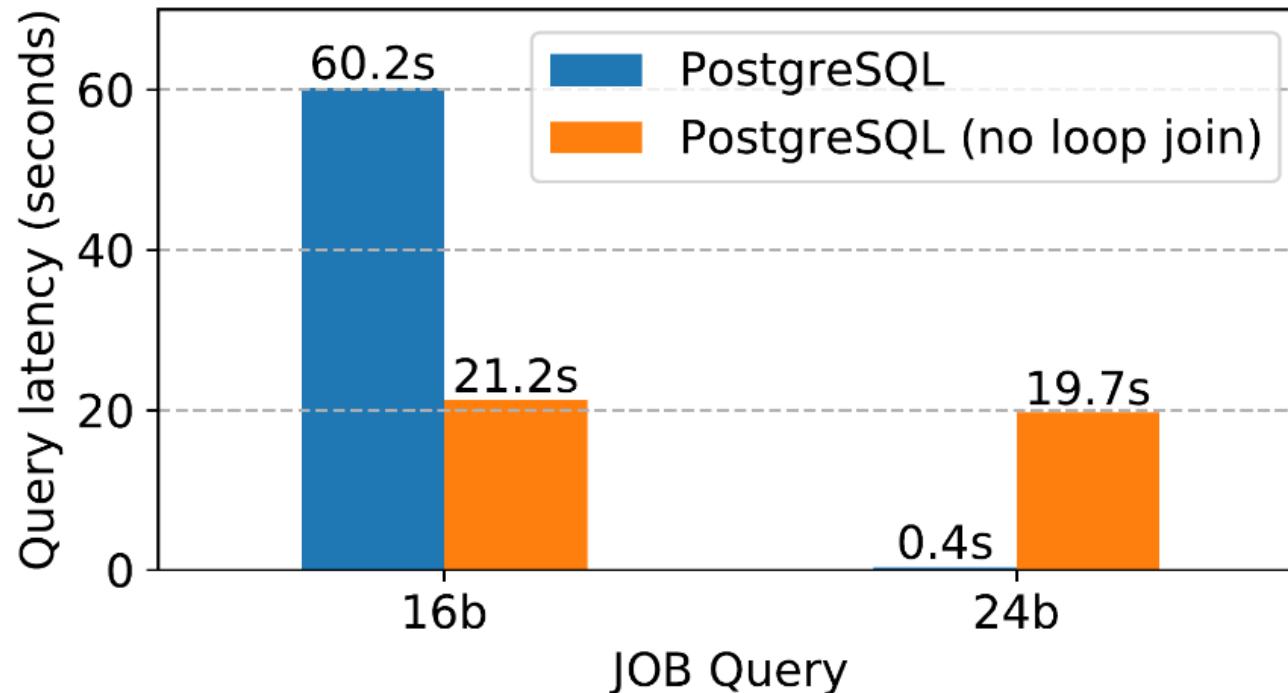


How to design a learned-based optimizer with
less training overhead and better robustness?



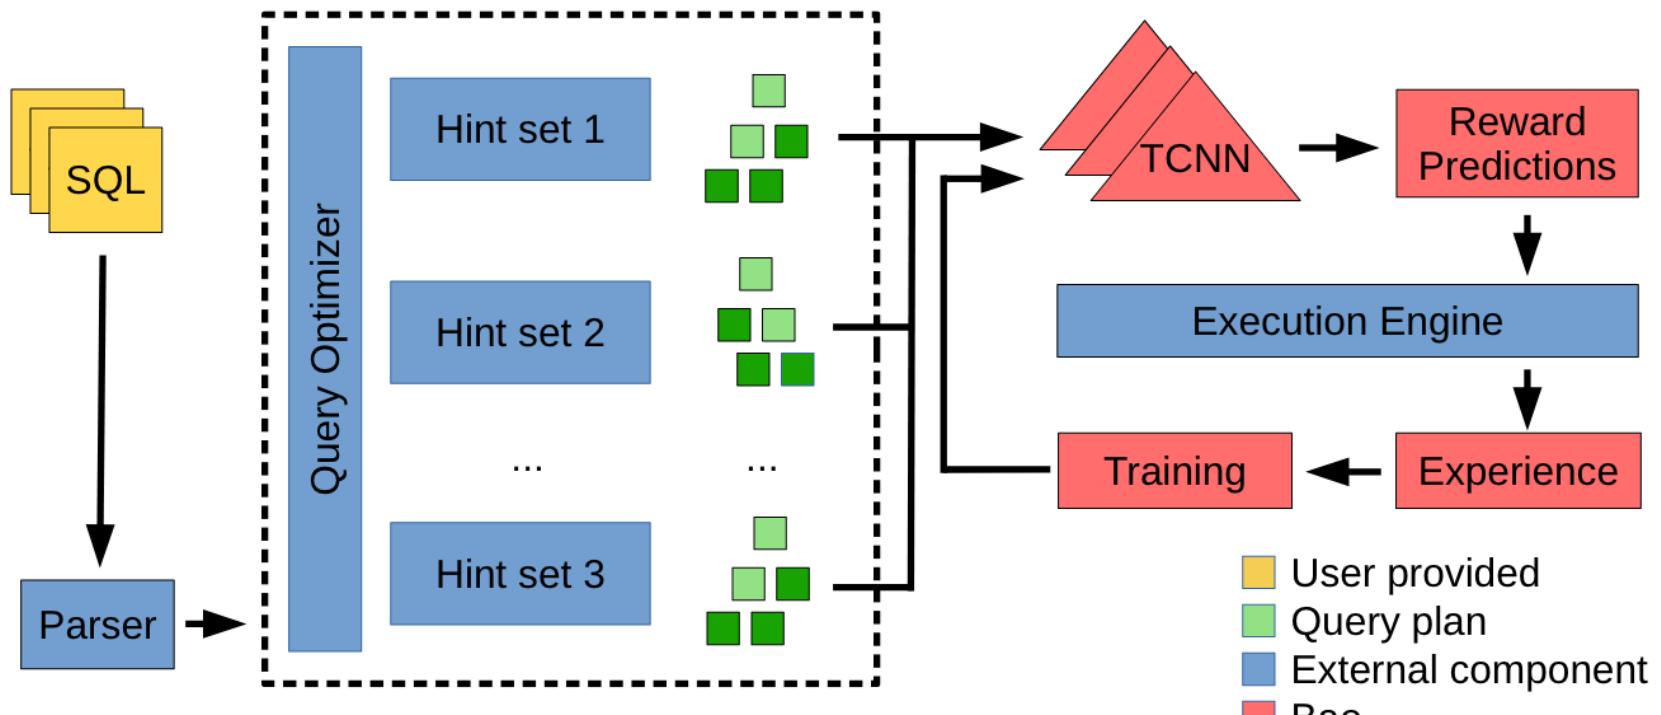
Architecture of learned optimizer enhancer

❖ Bao [1]



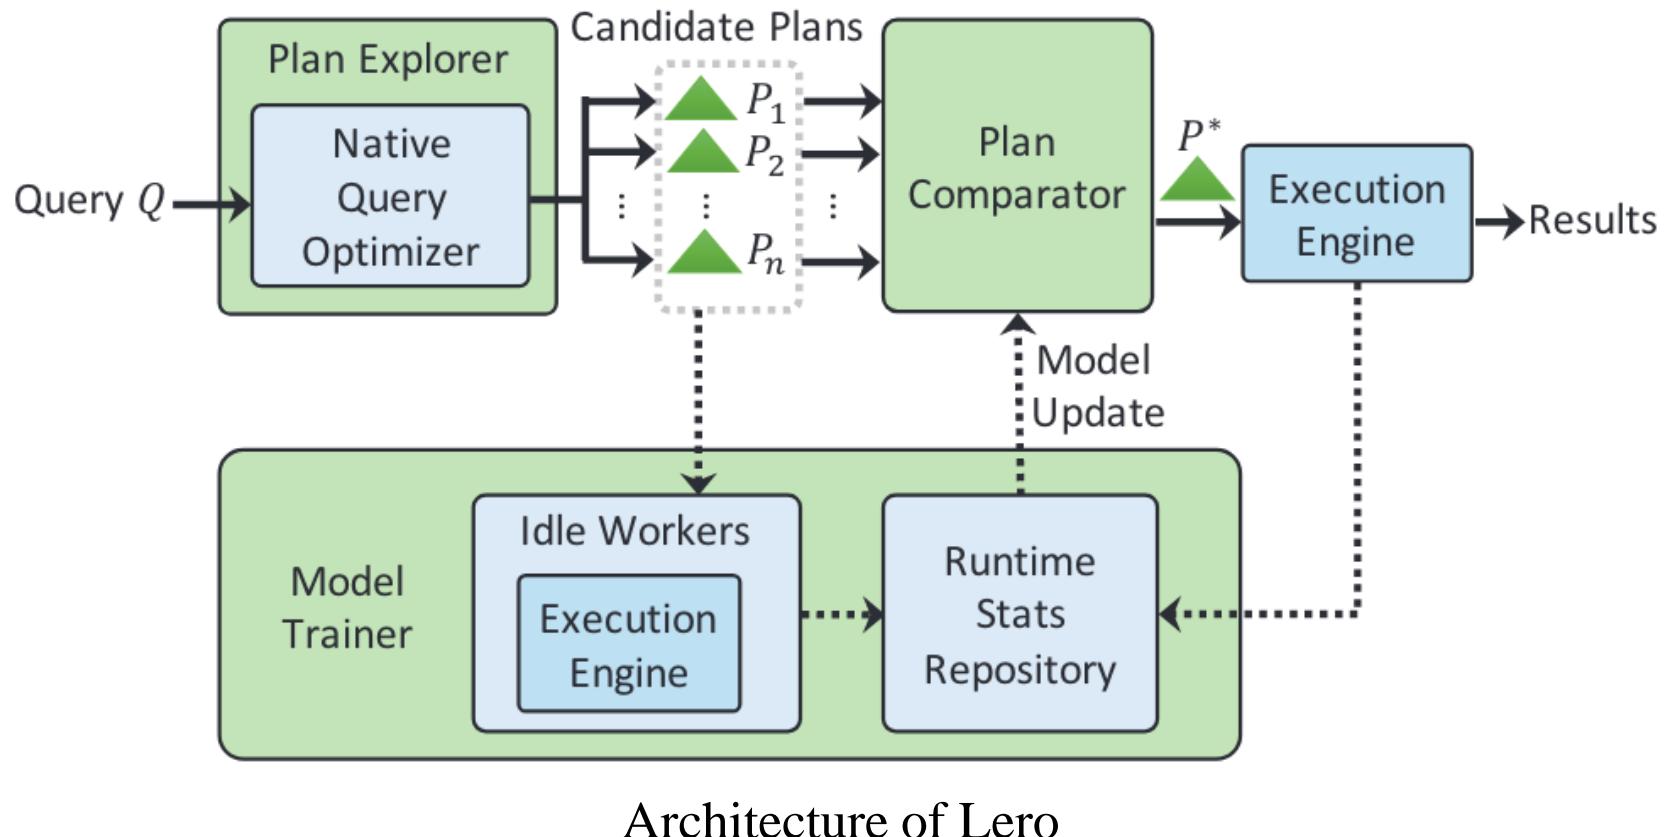
Disabling loop join improve (16b) but harm (24b) for JOB in PostgreSQL

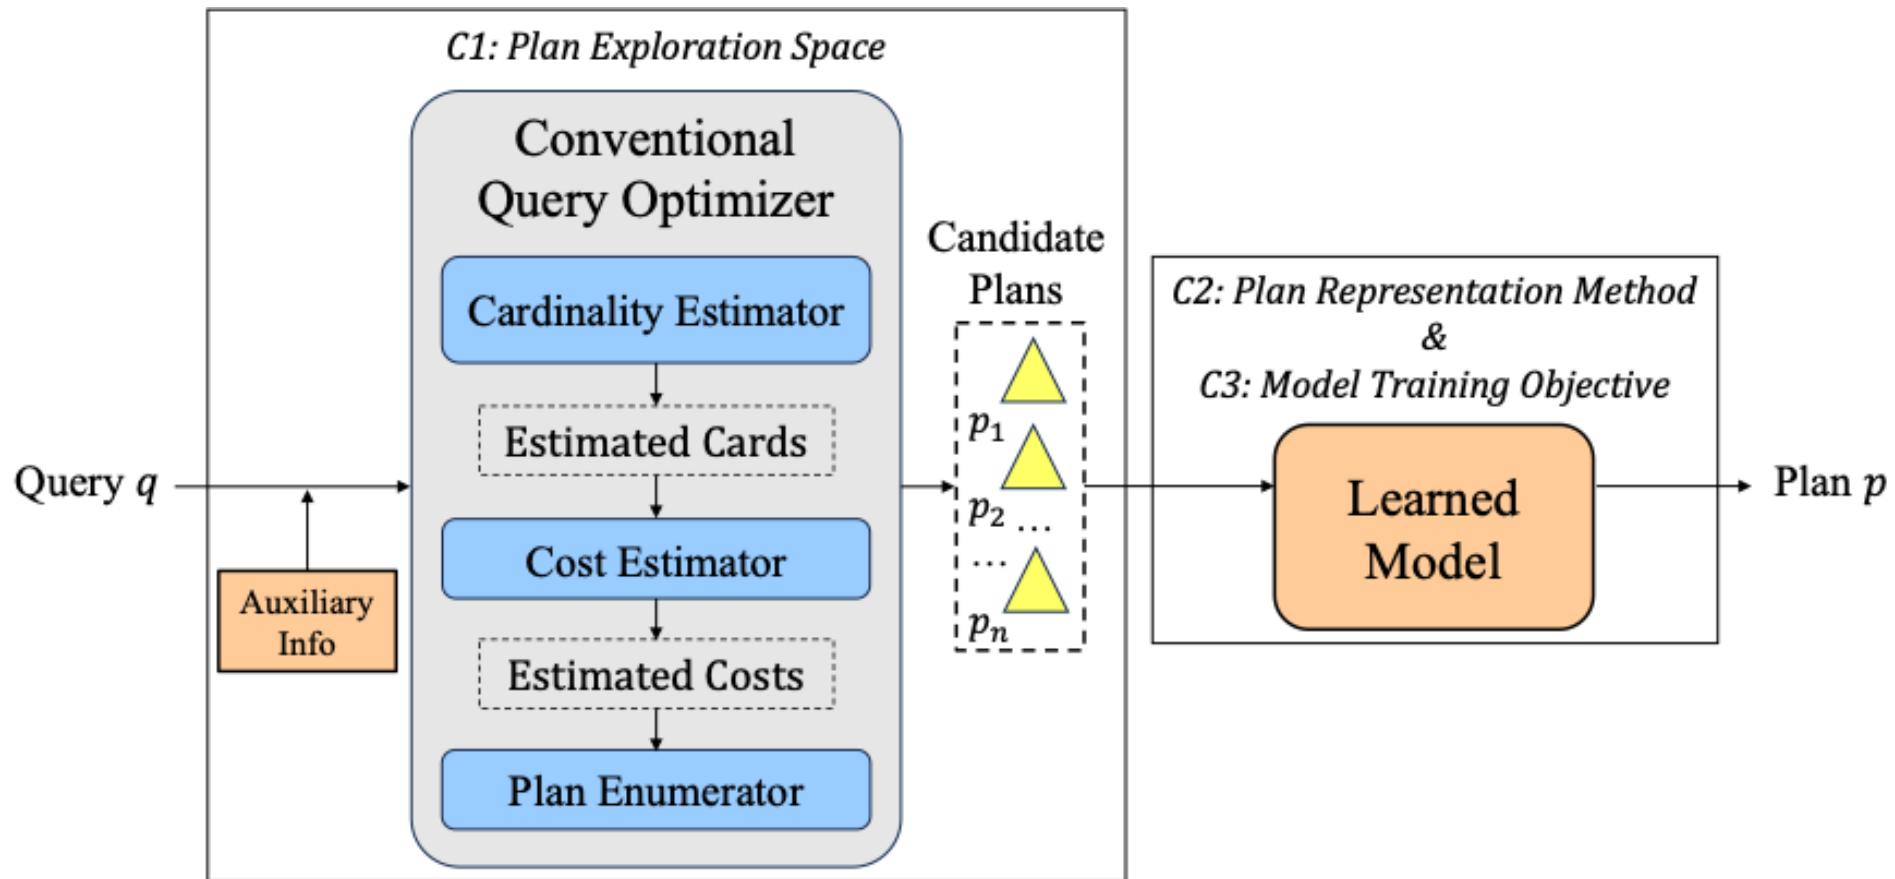
❖ Bao [1]



Architecture of Bao

❖ Lero [1]





Architecture of learned optimizer enhancer



Learned Optimizer Enhancer



❖ C1: Plan Exploration Space

❖ Existing methods

- Bao: variant hint assignment (e.g., enable hash join, enable index scan).
- Lero: cardinality perturbation (e.g., some cardinalities * 10 or * 0.1).

❖ Limitations of existing methods

- Bao: Fails when different query components demand distinct hints.
- Lero: Limited to local plan space exploration.
- Neither method directly explores diverse join order.

❖ C2: Plan Representation Method

❖ Existing methods

- Both Bao and Lero adopt Tree-CNN

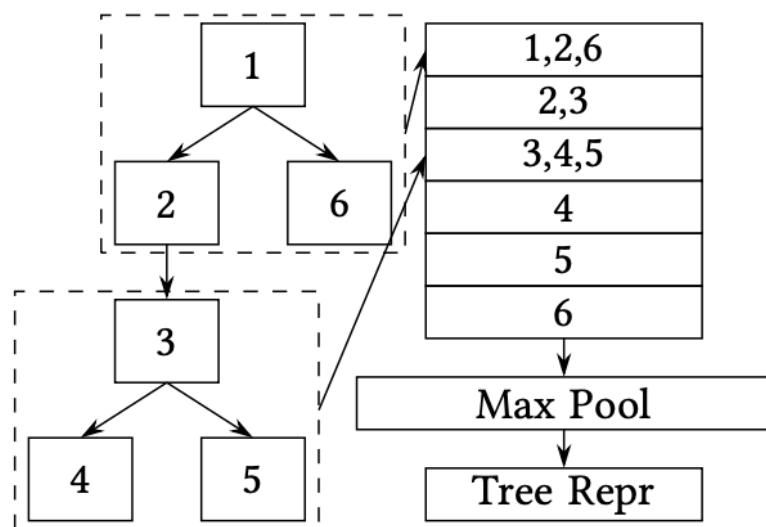
❖ Limitations of existing methods

- Loss of global structure information:

The max pooling-based feature

aggregation in Tree-CNN fails to

preserve hierarchical dependencies between nodes.



Tree-CNN procedure



Learned Optimizer Enhancer



❖ C3: Model Training Objective

❖ Existing methods

- Bao: learn to predict the execution time, taking MSE the loss function

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2,$$

\hat{y}_i is the predicted time and y_i is the actual time.

- Lero: learn to rank the plans of query, taking cross-entropy loss function

$$l = -\log(\sigma(\hat{y}_i - \hat{y}_j)),$$

\hat{y}_i and \hat{y}_j are the predicted time plan i and j, σ is sigmoid function.

This loss function means that plan i has larger actual time than plan j.

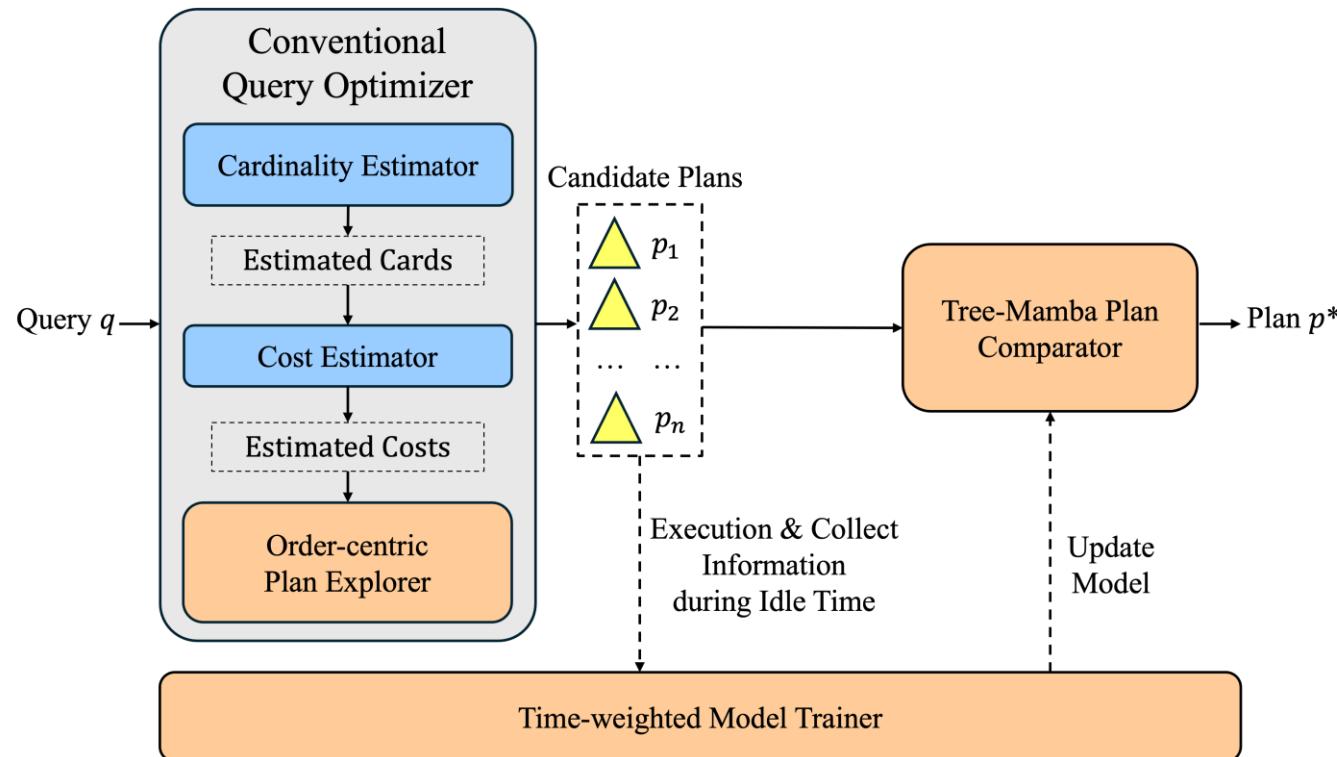


❖ C3: Model Training Objective

❖ Limitations of existing methods

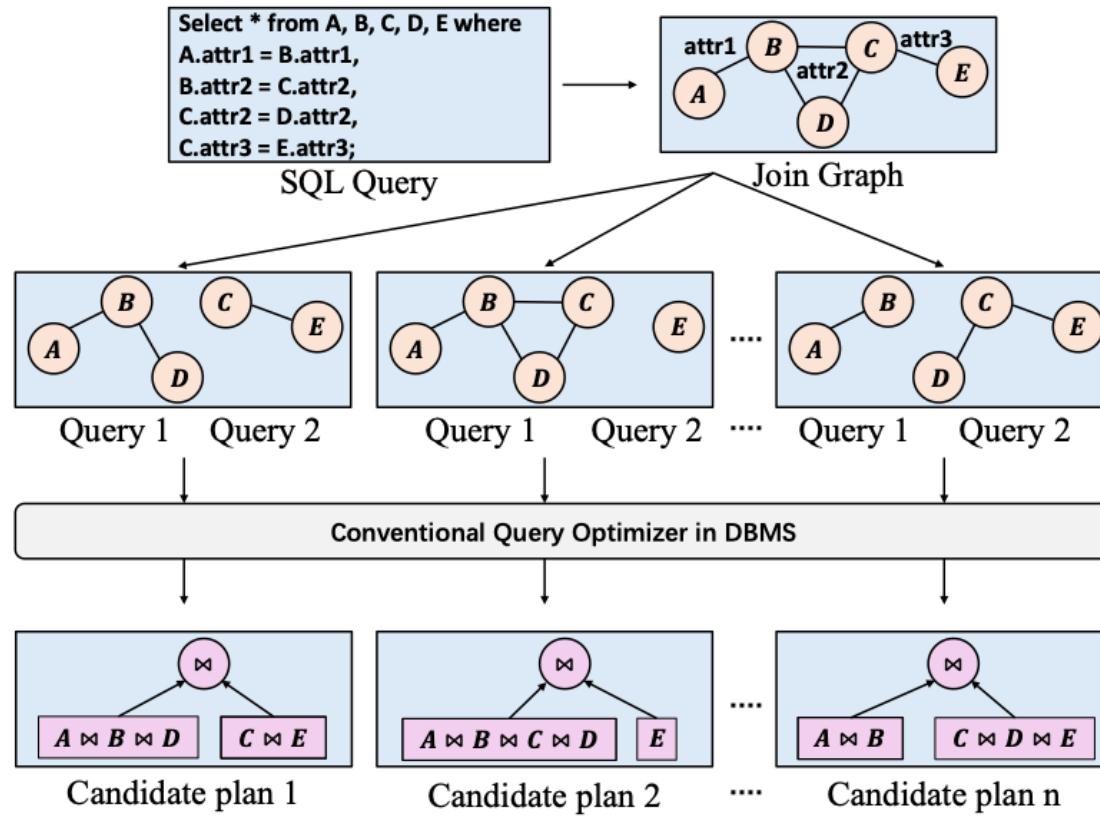
- Bao: learn to predict the execution time is impractical due to inherent estimation errors in query optimization.
- Lero: learn-to-rank enhances the model training but neglects sample importance differentiation
 - For example, two query pairs with execution time differences of 0.1s vs. 10s are treated equally by Lero's loss function, despite the latter being significantly more critical for performance optimization

❖ Athena [1]



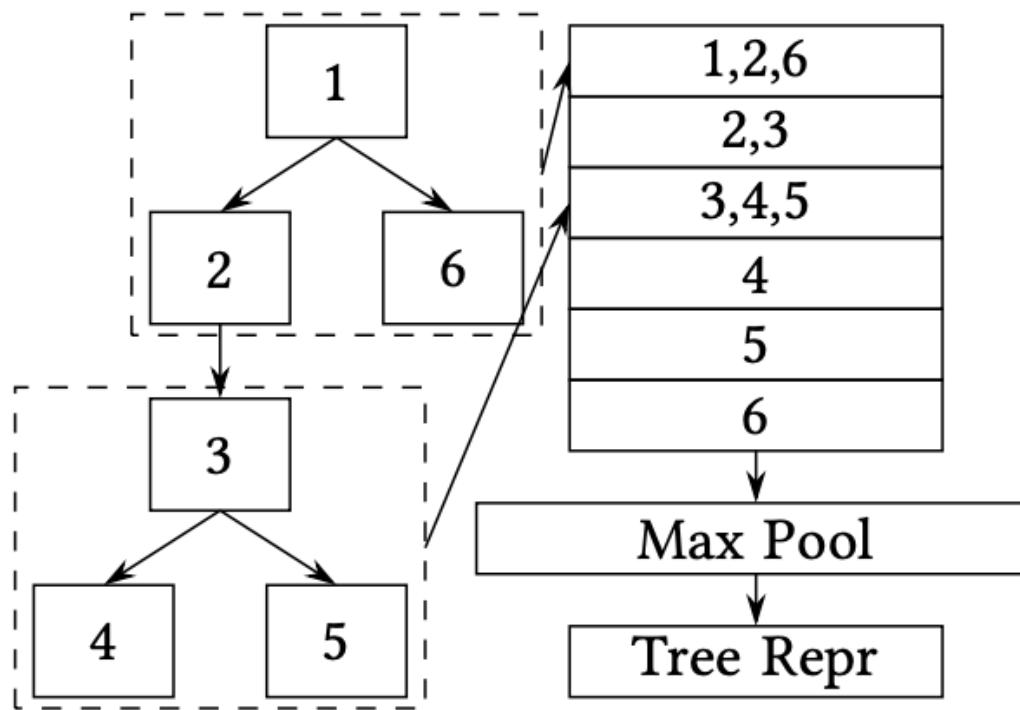
Architecture of Athena

❖ Athena [1] - Order-centric Plan Explorer

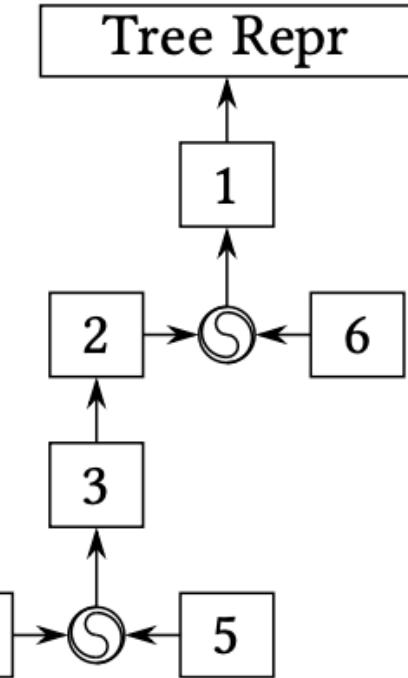


Plan explorer of Athena

❖ Athena [1] - Tree-Mamba Plan Comparator

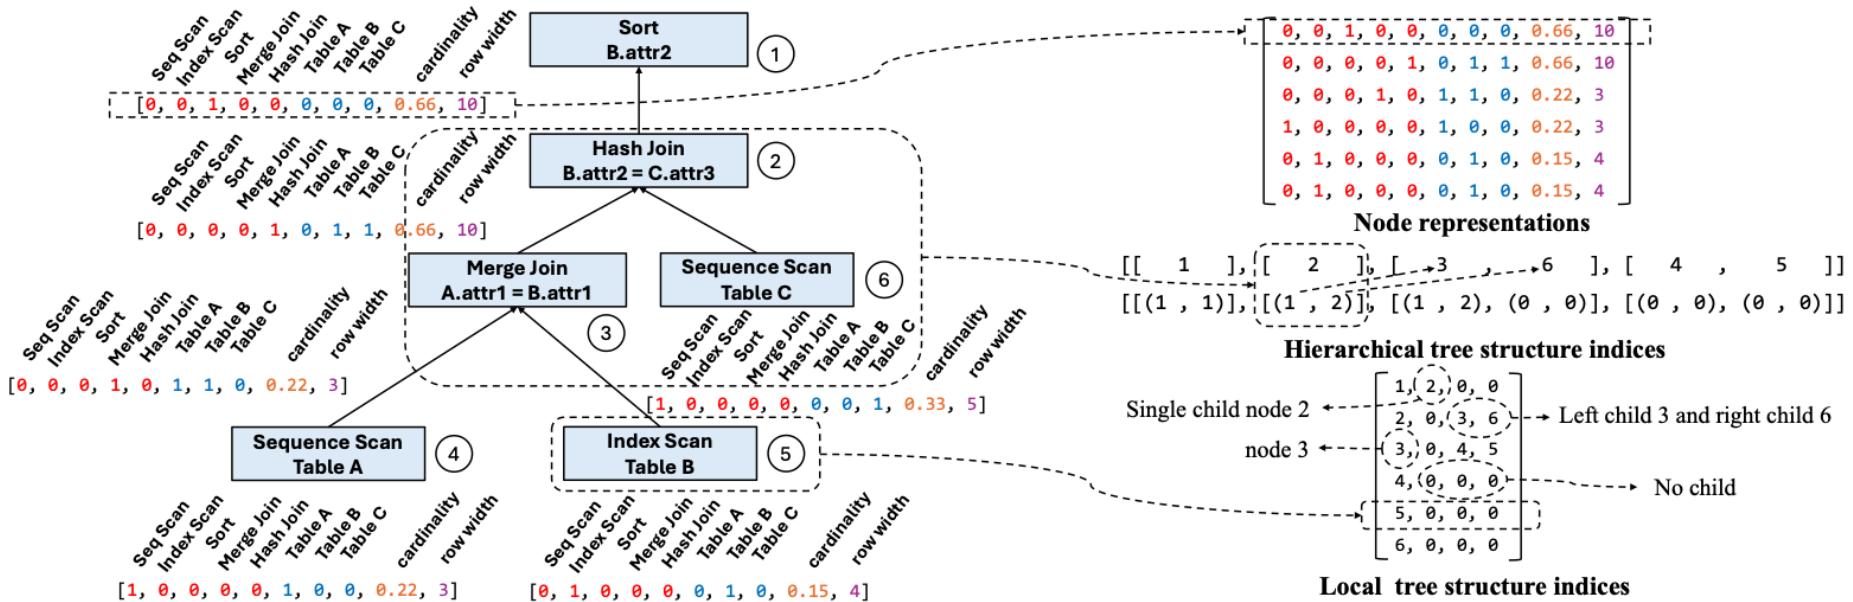


(a) Tree-CNN procedure



(b) Tree-RNN procedure

❖ Athena [1] - Tree-Mamba Plan Comparator - Model

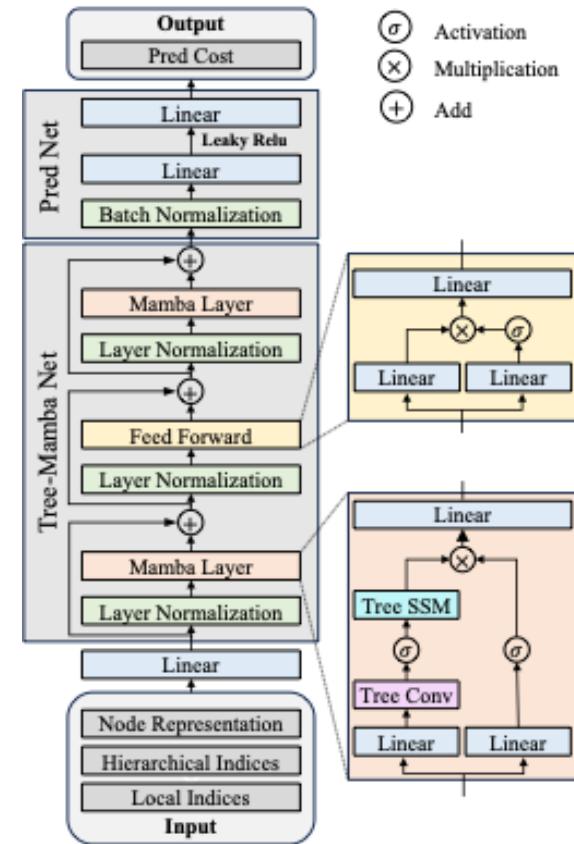


Featurizer of Athena

❖ Athena [1] - Tree-Mamba Plan Comparator - Model

❖ New design

- New Tree Conv could distinguish nodes with one or two child nodes.
- Adopt Mamba to represent the tree structure data with larger hidden space
- Optimized matrix multiplication.



Model design of Athena



Learned Optimizer Enhancer



❖ Athena [1] - Time-weighted Model Trainer

Pair time difference weight
Distinguish the importance of
different pair sample

$$\mathcal{L}'(\mathcal{B}; \tau, \gamma) = -\frac{\sum_{i=1}^n w_i^{\frac{1}{\tau}} \text{TaILr}(P(p_1 \triangleright p_2); \gamma)}{\sum_{i=1}^n w_i^{\frac{1}{\tau}}}$$

$\text{TaILr}(P; \gamma) = \frac{\log(\gamma + (1 - \gamma)P)}{1 - \gamma}$

TaILr loss function
Avoid overconfidence for low loss sample



Learned Optimizer Enhancer



❖ Athena [1]

The end-to-end latency (in seconds) on 4 benchmarks

Workload	JOB	STATS-CEB	TPC-DS	DSB
PostgreSQL	1666	8181	3773	1979
Bao	>40611	>14605	2089	2399
Lero	1661	7841	1106	1649
Athena	954	4185	663	722



Other AI4DB Topics



- ❖ Self-tuning Database Configuration
- ❖ Automated Index Recommendation
- ❖ Learned Index
- ❖ Security Enhancement and Anomaly Detection
- ❖ Automated Operations and Resource Scheduling
- ❖ ...



谢谢！

DBGroup @ SUSTech

Dr. Bo Tang (唐博)

tangb3@sustech.edu.cn

