



Transaction Processing

南方科技大学

唐 博

tangb3@sustech.edu.cn





Outline



- ❖ Basic concepts on transaction processing
- ❖ What is a serializable schedule?
- ❖ How to test for a serializable schedule?
- ❖ What is a recoverable schedule?

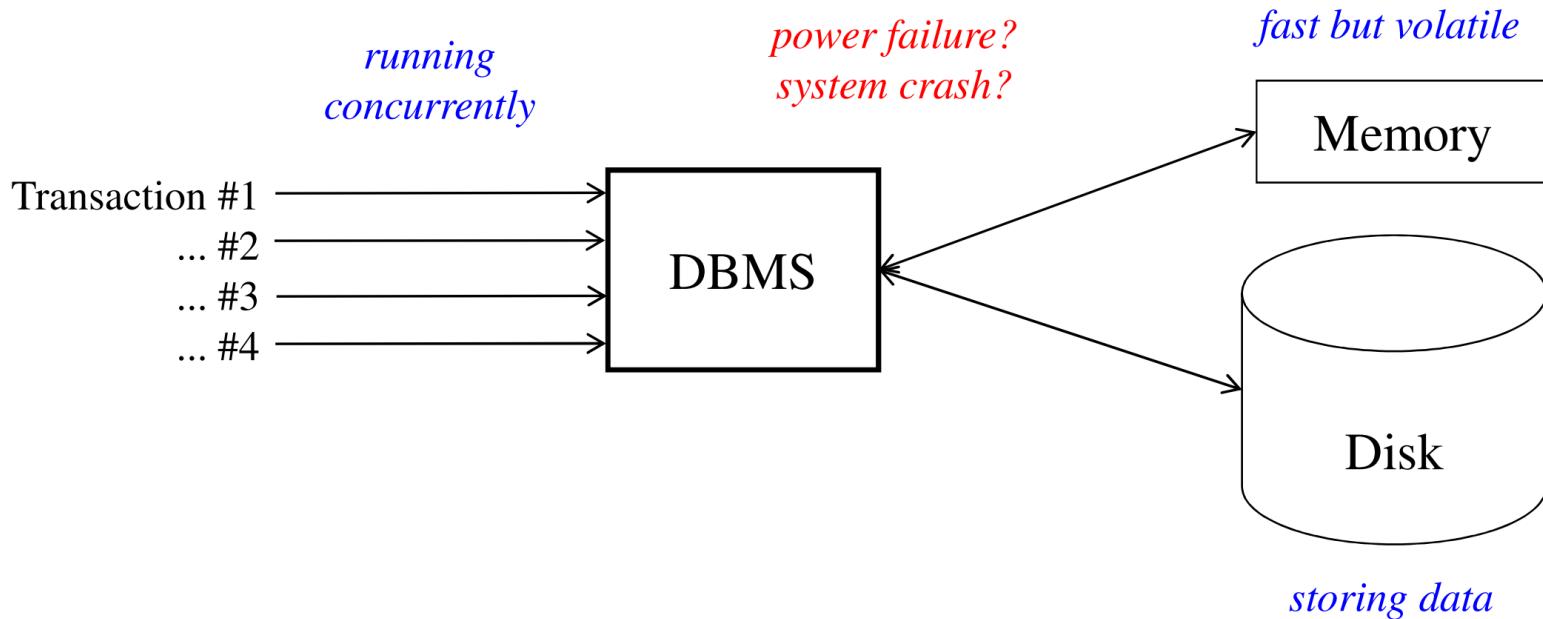


Outline



- ✓ Basic concepts on transaction processing
- ❖ What is a serializable schedule?
- ❖ How to test for a serializable schedule?
- ❖ What is a recoverable schedule?

Transaction Processing



- ❖ What is a transaction?
- ❖ Which properties should a DBMS guarantee?

ACID properties

(Atomicity, Consistency, Isolation, Durability)



What is a Transaction?



- ❖ **Transaction:** a unit of program that accesses and updates data items

Transaction to transfer \$v
from account A to account B:

1. **read**(A)
2. $A := A - v$
3. **write**(A)
4. **read**(B)
5. $B := B + v$
6. **write**(B)

❖ Operations

- ❖ read (X), write (X) (our focus)
- ❖ arithmetic operations (we do not care)
- ❖ commit or abort at end of transaction (to be discussed later)



Transaction in SQL



- ❖ SQL extensions to support transactions
- ❖ E.g., Transact-SQL (Microsoft, Sybase), PL/SQL (Oracle)
 - ❖ Can declare variables
 - ❖ Support flow control statements (e.g., if, else, while)
 - ❖ Support multiple SQL statements (e.g., SELECT, INSERT, DELETE, UPDATE)
 - ❖ COMMIT and ROLLBACK statements
 - ❖ <https://en.wikipedia.org/wiki/Transact-SQL>
- ❖ For simplicity, we write transactions like on the previous page



ACID properties: Consistency



- ❖ **Consistency:** Each single transaction preserves the consistency of the database
 - ❖ The consistency requirement depends on the application
 - ❖ Example: a fund transfer transaction should not create / destroy money
 - ❖ [Manual] Application programmers should avoid bugs in transactions
 - ❖ [Automatic] Check consistency at runtime by using integrity constraints
 - ❖ But expensive!

A *buggy* transaction to transfer \$x from account A to account B:

1. **read(A)**
2. $A := A - x$
3. **write(A)**
4. **read(B)**
5. $B := B + 2x$
6. **write(B)**



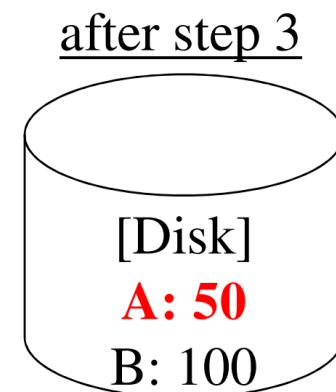
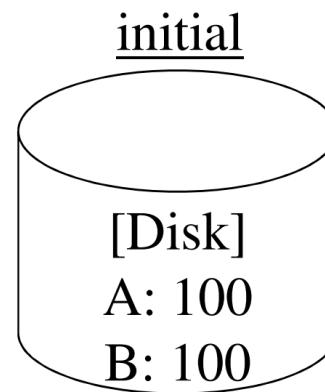
ACID properties: Atomicity

The logo for DBGroup is in the top right corner, featuring a red stylized flame icon above the text "DBGroup".

- ❖ **Atomicity:** Either all operations or none operations of the transaction are reflected in the database
 - ❖ Suppose that the system updates data items in the disk immediately
 - ❖ What happens if the system **fails** at this point?
 - ❖ E.g., power failure

Transaction to transfer \$50 from account A to account B:

1. **read(A)**
2. $A := A - 50$
3. **write(A)**
4. **read(B)**
5. $B := B + 50$
6. **write(B)**





ACID properties: Durability

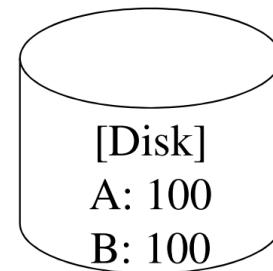


- ❖ **Durability:** After a transaction completes successfully, the changes made to the database must persist, even if there are system failures
 - ❖ Suppose that the system caches data items in the main memory
 - ❖ After the transaction has completed, what happens if the system **fails** at this point?

Transaction to transfer \$50 from account A to account B:
1. **read(A)**
2. $A := A - 50$
3. **write(A)**
4. **read(B)**
5. $B := B + 50$
6. **write(B)**

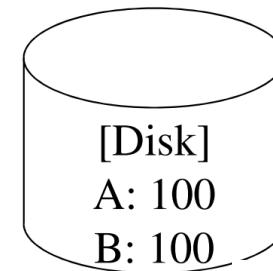
after step 6

[Memory]
A: 50
B: 150



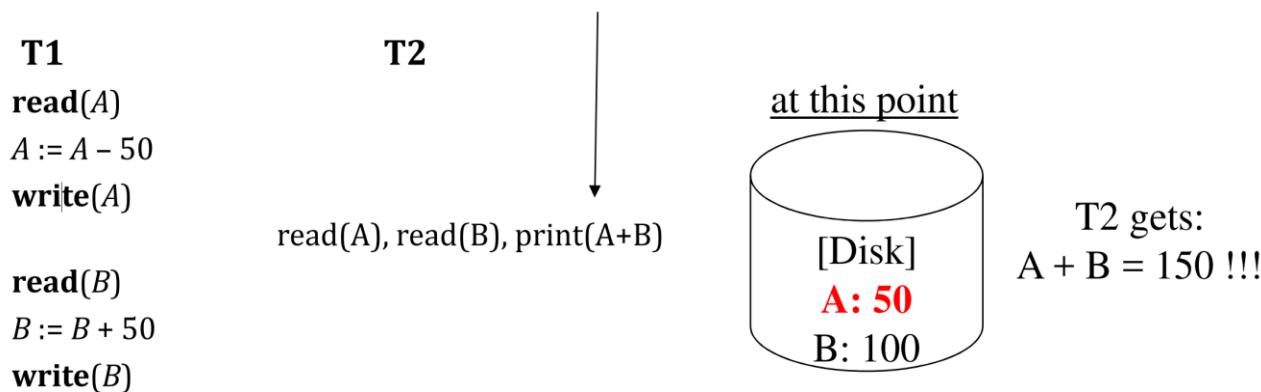
power failure

[Memory]

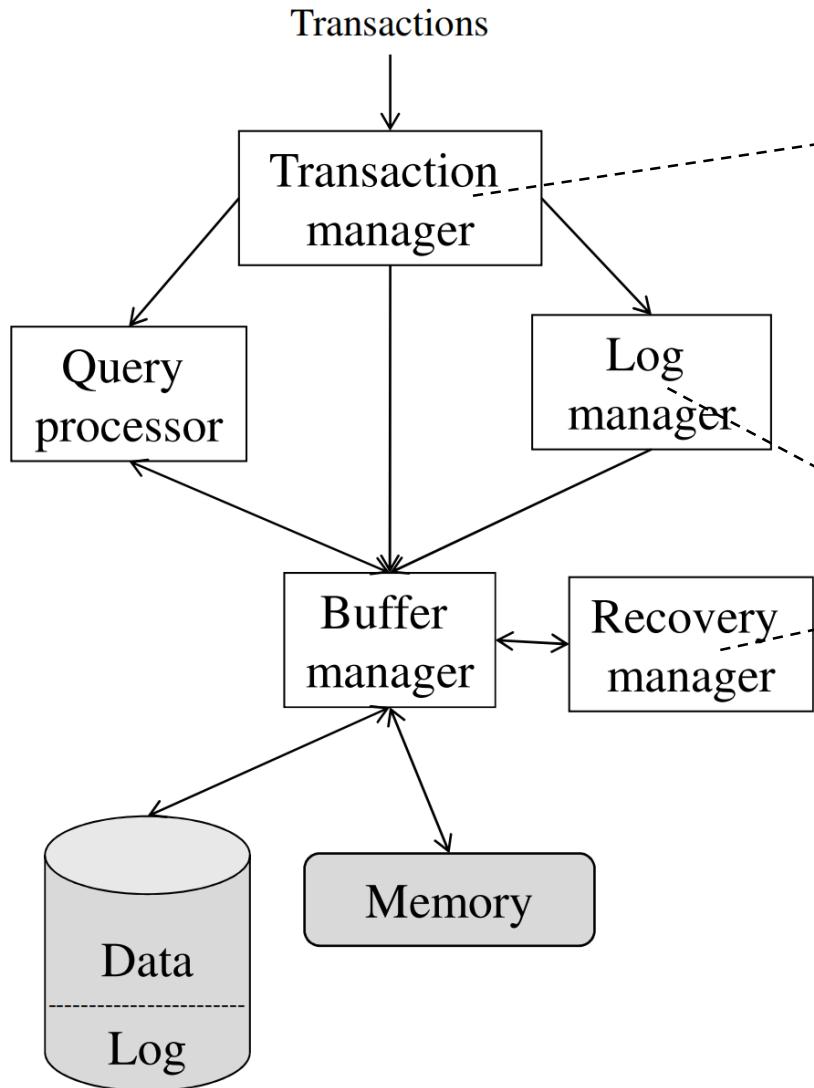


ACID properties: Isolation

- ❖ Advantages for DBMS to run multiple transactions concurrently
 - ❖ Better utilization of CPU and disk
 - ❖ a transaction is using the CPU while another is using the disk
 - ❖ Reduced average response time for transactions:
 - ❖ short transactions need not wait behind long ones
- ❖ **Isolation:** Each transaction cannot see any intermediate transaction result from concurrently executing transactions
 - ❖ Example: Transaction T2 can access the partially updated database (caused by T1) → T2 may see an inconsistent database



DBMS components for transaction processing



- ❖ Coordinates transactions that operate in parallel and access shared data (‘*I*’)
 - ❖ concurrency control schemes
- ❖ Ensure that failures do not corrupt persistent data (‘*A*’ , ‘*D*’)
 - ❖ logging and recovery



Transaction Operations



- ❖ Data access operations: *read, write*
 - ❖ Advanced operations (e.g., *increment* and *decrement*) can be executed as *atomic* actions
- ❖ Transaction operations: *Start, Commit, Abort*
 - ❖ **Start:** begins a new transaction
 - ❖ **Commit:** terminates a transaction successfully, and all of its effects should be made permanent
 - ❖ **Abort:** terminates a transaction abnormally and all of its effects should be eliminated
 - ❖ We roll back an aborted transaction and undo its effects



Commit/Abort Example



Transfer

```
Start;  
input(from-ac, to-ac, amount);  
temp = read(Account[from-ac]);  
if (temp < amount) {  
    output("insufficient fund");  
    abort; ← When a transaction  
fails to complete its  
execution  
} else {  
    write(Account[from-ac], temp - amount);  
    temp = read(Account[to-ac]);  
    write(Account[to-ac], temp + amount);  
    commit; ← When a transaction  
successfully completes  
its execution  
    output("transfer completed");  
}  
// Note: we have omitted other error checking in this example, e.g., input check.
```

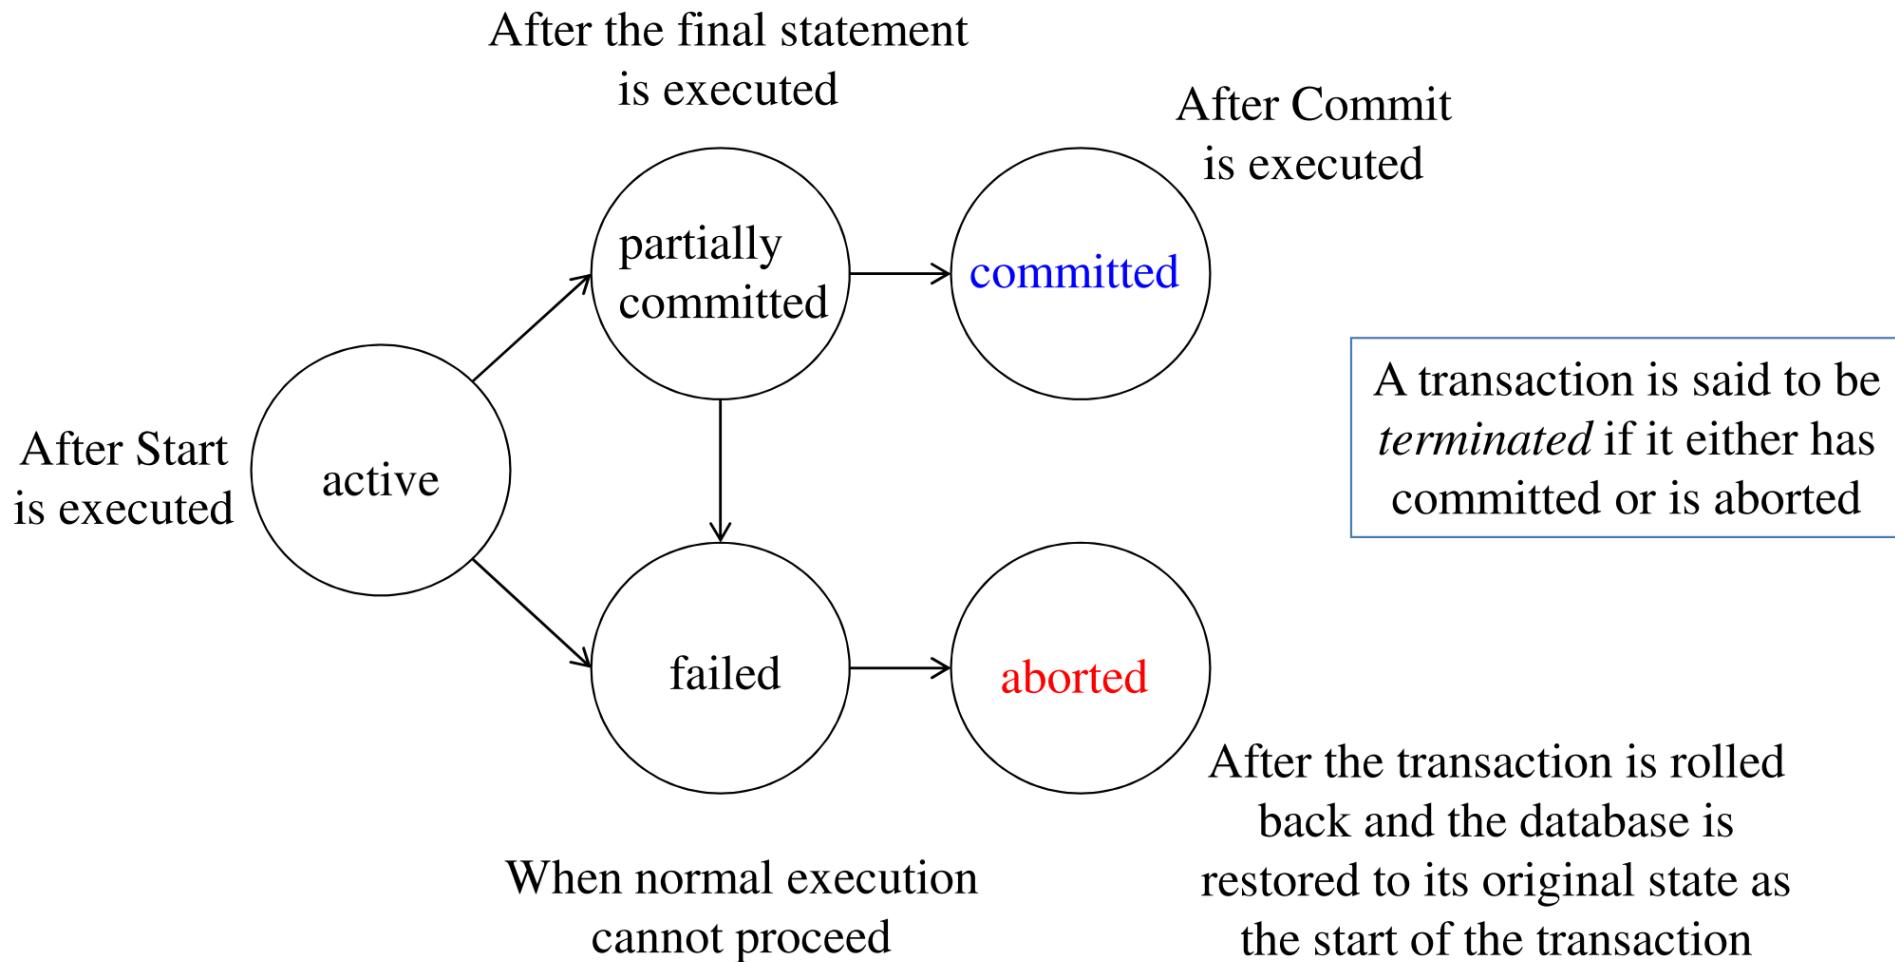


More on Abort and Commit



- ❖ Three reasons to abort a transaction:
 - ❖ Its own logic (e.g., insufficient fund in “from-ac”)
 - ❖ Caused by DBMS due to system failures
 - ❖ Caused by DBMS due to concurrency control (e.g., the transaction reads an uncommitted value)
- ❖ To ensure **atomicity**, we must **undo** any changes caused by the aborted transaction
 - ❖ Also, the effect of an aborted transaction must not be seen by other transactions
- ❖ Executing a transaction’s **commit** guarantees that
 - ❖ The transaction’s effects will not be undone in future and
 - ❖ Such effects will survive the system’s failures in future

Transaction States





Aborting a Transaction



- ❖ Two ways to deal with an aborted transaction:
 - ❖ *Restart*: if the abort is due to hardware failure or concurrency control, the system could restart the transaction after the system recovers
 - ❖ *Kill*: if the abort is due to logical error, the transaction is discarded



Outline



- ✓ Basic concepts on transaction processing
- ✓ What is a serializable schedule?
- ❖ How to test for a serializable schedule?
- ❖ What is a recoverable schedule?

- ❖ **Schedule** – it specifies the execution order of instructions of concurrent transactions
 - ❖ must contain all instructions of those transactions
 - ❖ must preserve the order of instructions within each individual transaction

| <u>Transactions</u> | | <u>a schedule</u> | | <u>not a schedule!</u> | |
|---------------------|----------|-------------------|----------|------------------------|---------|
| T_1 | T_2 | T_1 | T_2 | T_1 | T_2 |
| read(A) | read(A) | read(A) | | write(A) | |
| write(A) | write(A) | | write(A) | read(A) | read(A) |

↓
Time

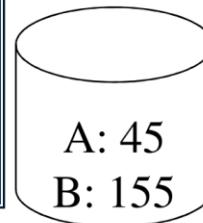
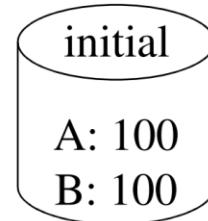
Schedules

- ❖ **Serial schedule:** the next transaction begins only after the current one completes
- ❖ Let T1 transfer \$50 from A to B, and T2 transfer 10% of the balance from A to B

Serial schedule 1

$\langle T_1, T_2 \rangle$

| T_1 | T_2 |
|--|---|
| $\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ | $\text{read}(A)$ $temp := A * 0.1$ $A := A - temp$ $\text{write}(A)$ $\text{read}(B)$ $B := B + temp$ $\text{write}(B)$ |



Serial schedule 2

$\langle T_2, T_1 \rangle$

| T_1 | T_2 |
|-------|---|
| | $\text{read}(A)$ $temp := A * 0.1$ $A := A - temp$ $\text{write}(A)$ $\text{read}(B)$ $B := B + temp$ $\text{write}(B)$ |

| | T_1 |
|--|--|
| | $\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ |



Examples of Concurrent Schedules



Schedule 3

| T ₁ | T ₂ |
|--------------------------------------|---|
| read(A) $A := A - 50$ write(A) | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) |
| read(B) $B := B + 50$ write(B) | read(B) $B := B + temp$ write(B) |

Schedule 4

| T ₁ | T ₂ |
|--|--|
| read(A) $A := A - 50$ | read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) |
| write(A) read(B) $B := B + 50$ write(B) | $B := B + temp$ write(B) |

A, B
100, 100

90, 100

50, 100

- ❖ Schedule 3 is not a serial schedule, but it is equivalent to Schedule 1
- ❖ The sum $A + B$ is preserved in Schedules 1, 2, 3, but it is not preserved in Schedule 4



Serializability



- ❖ **Basic Assumption** – Each transaction preserves database consistency
 - ❖ Thus any serial schedule preserves database consistency
- ❖ A schedule is **serializable** if it is equivalent to a serial schedule. Two forms of schedule equivalence:
 1. **conflict serializability**
 2. **view serializability**
- ❖ *Simplified view of transactions*
 - ❖ We ignore operations other than **read** and **write**
 - ❖ Transactions may perform any computation on data in local buffers (in main memory) in between reads and writes



Conflicting Instructions



| | $l_i = \text{read}(Q)$ | $l_i = \text{write}(Q)$ |
|-------------------------|------------------------|-------------------------|
| $l_j = \text{read}(Q)$ | No conflict | Conflict |
| $l_j = \text{write}(Q)$ | Conflict | Conflict |

- ❖ Let l_i and l_j be instructions of transactions T_i and T_j respectively
- ❖ l_i and l_j are **conflict** when
 - ❖ both of them access the same data item Q_j , and
 - ❖ at least one of them is **write**(Q)
- ❖ The result remains the same if we swap two consecutive instructions l_i and l_j that do not conflict
- ❖ A conflict between l_i and l_j forces a temporal order between them



Conflict Serializability



- ❖ Two schedules S and S' are **conflict equivalent** if S can be transformed into S' by a series of swapping non-conflicting instructions
- ❖ A schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule.

Any conflict
between these
two instructions?

| T_1 | T_2 |
|-----------------------------|--|
| read(A) write(A) | |
| read(B) write(B) | read(A) write(A) read(B) write(B) |

Schedule S

| T_1 | T_2 |
|--|--|
| read(A) write(A) read(B) write(B) | read(A) write(A) read(B) write(B) |

A serial schedule



Conflict Serializability (Cont.)



❖ Is this schedule conflict serializable?

Can we swap them ?

| T_3 | T_4 |
|--------------|--------------|
| read(Q) | |
| write(Q) | write(Q) |

Can we swap them ?

❖ NO because

- ❖ Possible serial schedules are: $< T_3, T_4 >$ or $< T_4, T_3 >$
- ❖ But we are unable to swap instructions in the above schedule to obtain either serial schedule



View Serializability



- ❖ Let S and S' be two schedules with the same set of transactions. S and S' are **view equivalent** if the following three conditions are satisfied, for each data item Q ,
 1. The transaction (if any) that performs the first **read(Q)** operation in schedule S must also perform the first **read(Q)** operation in schedule S' .
 2. The transaction (if any) that performs the final **write(Q)** operation in schedule S must also perform the final **write(Q)** operation in schedule S' .
 3. If in schedule S transaction T_i reads the value of Q produced by transaction T_j , then in schedule S' also transaction T_i must read the value of Q produced by the same **write(Q)** operation of transaction T_j .
- ❖ View equivalence is also based purely on **reads** and **writes** alone.
- ❖ A schedule S is **view serializable** if it is view equivalent to a serial schedule.

View Serializability (Cont.)

- ❖ Every conflict serializable schedule is also view serializable.
- ❖ Some schedule (e.g., the one below) is view-serializable but not conflict serializable.
 - ❖ Such schedule has **blind write(s)**

| T_3 | T_4 | T_6 |
|--------------|--------------|--------------|
| read(Q) | | |
| write(Q) | write(Q) | write(Q) |

- ❖ Which serial schedule is the above equivalent to?
 - ❖ $< T_3, T_4, T_6 >$ or $< T_4, T_3, T_6 >$



Serializability of Schedules



All schedules

Schedules equivalent to some serial schedule

View serializable schedules

Conflict serializable schedules

Serial schedules



Outline



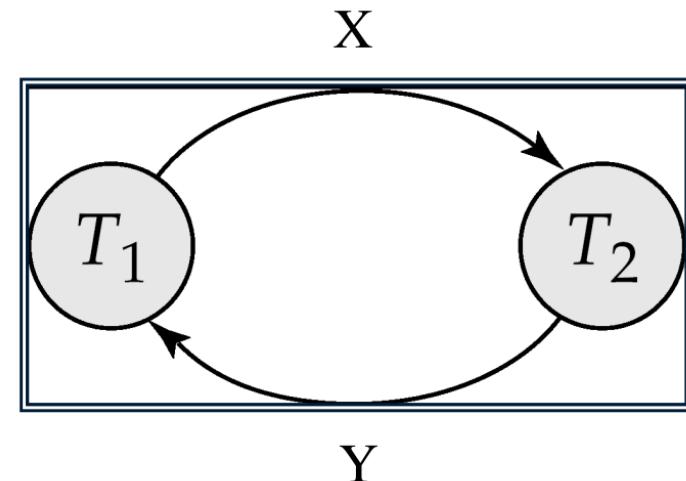
- ✓ Basic concepts on transaction processing
- ✓ What is a serializable schedule?
- ✓ How to test for a serializable schedule?
- ❖ What is a recoverable schedule?

❖ Consider a schedule of a set of transactions T_1, T_2, \dots, T_n

❖ Precedence graph

- ❖ A *vertex* denotes a transaction (name)
- ❖ A *directed arc* from T_i to T_j if the two transactions conflict on a data item, and T_i accessed that data item earlier
- ❖ Label the arc by the corresponding item

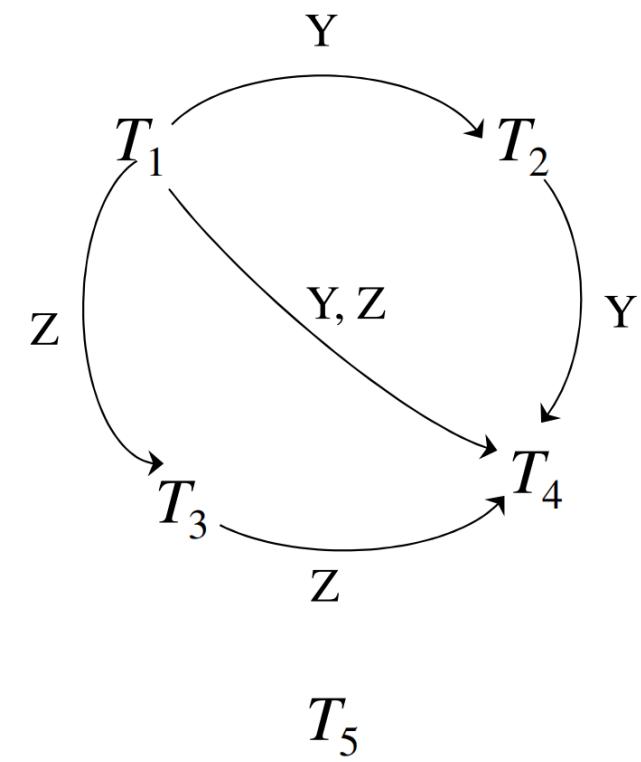
| T_1 | T_2 |
|--------------|-----------------------------|
| write(X) | |
| read(Y) | read(X) write(Y) |



Example: Precedence Graph



| T_1 | T_2 | T_3 | T_4 | T_5 |
|---------------------|---------------------|----------|--|-------------------------------|
| read(Y) read(Z) | read(X) | | | |
| | read(Y) write(Y) | | | read(V) read(W) read(W) |
| read(U) | | write(Z) | read(Y) write(Y) read(Z) write(Z) | |
| read(U) write(U) | | | | T_5 |



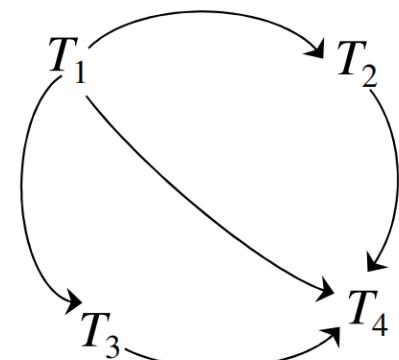


Test for Conflict Serializability



- ❖ A schedule is conflict serializable if and only if its precedence graph is **acyclic**
- ❖ If precedence graph is acyclic, we can obtain the **serializability order** by a *topological sorting* of the graph
 - ❖ Iteratively pick a vertex with no incoming arcs, and then remove all arcs from that vertex
 - ❖ Example:
Consider this precedence graph.
Find a serializability order for this schedule.

$< \underline{\hspace{1cm}}, \underline{\hspace{1cm}}, \underline{\hspace{1cm}}, \underline{\hspace{1cm}}, \underline{\hspace{1cm}} >$





Test for View Serializability



- ❖ The test for view serializability is very expensive
 - ❖ E.g., check the schedule against all possible serial schedules
- ❖ No known algorithm to solve this problem efficiently
- ❖ Some subclasses of view serializability can be checked efficiently
 - ❖ These algorithms are out of the scope of this course



Outline



- ✓ Basic concepts on transaction processing
- ✓ What is a serializable schedule?
- ✓ How to test for a serializable schedule?
- ✓ What is a recoverable schedule?

We must consider the effect of transaction failures on concurrently running transactions.



Recoverable Schedules



- ❖ Recoverable schedule — if a transaction T_j reads a data item previously written by a transaction T_i , then the **commit** operation of T_i appears before the **commit** operation of T_j
- ❖ The following schedule is not recoverable if T_9 commits immediately after the read!

| T_8 | T_9 |
|----------|--------------------------|
| read(A) | |
| write(A) | |
| | read(A) commit |
| | read(B) |

- ❖ If T_8 may abort in future, T_9 would have read (and shown to the user) an inconsistent database state!
- ❖ Hence, DBMS must ensure that schedules are recoverable

Cascading Rollbacks

- ❖ **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks.
- ❖ Example: the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

| T_{10} | T_{11} | T_{12} |
|--|-----------------------------|-------------|
| read(A) read(B) write(A) | | |
| | read(A) write(A) | read(A) |

- ❖ If T_{10} fails, then we must also roll back T_{11} and T_{12}
- ❖ May undo a large amount of work 



Cascadeless Schedules



- ❖ **Cascadeless schedule** — cascading rollbacks cannot occur:
 - ❖ for each pair of transactions T_i and T_j such that
 T_j reads a data item previously written by T_i ,
the commit operation of T_i appears before the read operation of T_j
- ❖ Every cascadeless schedule is also recoverable
- ❖ It is desirable to restrict the schedules to cascadeless ones

| T_1 | T_2 | T_3 |
|---------------|--------------------------|---------------|
| read(A) | | |
| read(B) | | |
| write(A) | | |
| commit | | |
| | read(A) commit | |
| | | read(A) |
| | | read(A) |
| | | commit |



Recoverability of Schedules



All schedules

Recoverable schedules

Cascadless schedules

Serial schedules



Concurrency Control



- ❖ A database must provide a mechanism to ensure that all possible schedules are
 - ❖ either conflict or view **serializable**, and
 - ❖ are **recoverable** and preferably cascadeless
- ❖ A simple policy: execute only one transaction at a time.
 - ❖ This generates serial schedules, but provides poor concurrency.
 - ❖ Are serial schedules recoverable/cascadeless?
- ❖ Too late to test a schedule for serializability after executing it!
- ❖ **Goal** – develop concurrency control protocols (e.g., locking) to guarantee that schedules are serializable, recoverable and cascadeless
 - ❖ These protocols do not need to examine the precedence graph ☺



Weak Levels of Consistency



- ❖ Some applications accept weak levels of consistency and non-serializable schedules
 - ❖ E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - ❖ E.g., database statistics computed for query optimization can be approximate (why?)
 - ❖ Such transactions need not be serializable with respect to other transactions
 - ❖ Supported by setting the transaction isolation level:
<https://msdn.microsoft.com/en-us/library/ms173763.aspx>
- ❖ Tradeoff accuracy for performance



Conclusions



- ❖ ACID properties
- ❖ Isolation requirements for concurrent executions (I)
 - ❖ ensured by serializable schedules
 - ❖ conflict serializability, view serializability
 - ❖ [Future] we will study concurrency control schemes which produce serializable schedules
- ❖ Transaction failures (A,D)
 - ❖ ensured by recoverable and cascadeless schedules
 - ❖ [Future] we will study the recovery manager component which supports atomicity and durability



谢谢！

DBGroup @ SUSTech

Dr. Bo Tang (唐博)

tangb3@sustech.edu.cn

