



Concurrency Control


南方科技大学

唐 博

tangb3@sustech.edu.cn



Scheme



- ❖ Execute a concurrency control scheme
- ❖ Control the interaction among transactions
- ❖ Produce **serializable** schedule



If No Concurrency Control




Problems

- ❖ **Lost update**
 - ❖ A transaction's update is overwritten by another transaction
- ❖ **Dirty read**
 - ❖ A transaction reads data which was produced by uncommitted transaction
- ❖ **Non-repeatable read**
 - ❖ A transaction reads the same data item twice (or more), but obtains different values each time
- ❖ **Phantom read**
 - ❖ A transaction runs the same query twice (or more), but obtains different results each time
 - ❖ E.g., caused by INSERT, DELETE from other transactions



Lecture Objectives



- ❖ Two-Phase Locking
- ❖ Problems of Locking

- ❖ Timestamp Ordering
- ❖ Optimistic Concurrency Control





Industry



FYI: What are used in the industry?

- ❖ Microsoft SQL Server supports
 - ❖ Locking
 - ❖ Optimistic concurrency control
 - ❖ [https://technet.microsoft.com/en-us/library/ms189132\(v=sql.105\).aspx](https://technet.microsoft.com/en-us/library/ms189132(v=sql.105).aspx)
- ❖ Oracle Database supports
 - ❖ Locking
 - ❖ https://docs.oracle.com/cd/B19306_01/server.102/b14220/consist.htm#i5702



Idea of Locking



- ❖ Idea: mutual exclusion
 - ❖ Avoid transactions with conflicting operations to execute concurrently
 - ❖ Do you still remember “*conflicting operations*” ?

- ❖ Typically done by locking
 - ❖ Shared-lock (**lock-S**): a transaction must hold a lock-S on an item x before using $\text{Read}(x)$
 - ❖ Exclusive-lock (**lock-X**): a transaction must hold a lock-X on an item x before using $\text{Write}(x)$



Lock Conflict




- ❖ Two locks conflict if
 - ❖ they are on the *same item*,
 - ❖ issued by *different transactions*, and
 - ❖ at least one of them is an exclusive lock
- ❖ Compatibility matrix

	Lock-S	Lock-X
Lock-S	yes	no
Lock-X	no	no

Problem

Locking alone cannot avoid non-serializable schedule!



❖ We need a **locking protocol!**

- ❖ a set of rules for requesting & releasing locks in transactions
- ❖ to restrict the possible schedules



Basic Two-phase Locking



❖ The basic 2PL protocol

- ❖ A transaction T must hold a lock on an item x in the appropriate mode before accessing x
- ❖ T waits if a conflicting lock on x is being held by another transaction
- ❖ Once T releases a lock, it cannot obtain any other lock subsequently

❖ It divides a transaction into two phases:

- ❖ A *growing phase* (obtaining locks)
- ❖ A *shrinking phase* (releasing locks)

Why Two Phases?

Under 2PL,
this unlock
operation cannot
be placed here.

T_1	T_2
Lock-S(x) R(x) Unlock(x)	Lock-X(x) Lock-X(y) W(x) Unlock(x) W(y) Unlock(y)
Lock-S(y) R(y) Unlock(y)	

Under 2PL, this is
the earliest possible
time to unlock x

2PL makes sure that the conflicting operations of the two transactions are ordered the same way

Why Two Phases?

Red operations: have been executed

Black Operations: will consider in future

T_1	T_2
$\text{Lock-S}(x)$ $R(x)$	$\text{Lock-X}(x)$ $\text{Lock-X}(y)$ $W(x)$ $\text{Unlock}(x)$ $W(y)$ $\text{Unlock}(y)$
$\text{Lock-S}(y)$ $\text{Unlock}(x)$ $R(y)$ $\text{Unlock}(y)$	

T_2 lock x operation cannot be granted at this point

T_2 can lock x after T_1 unlocks x

Why Two Phases?

T_1	T_2
$\text{Lock-S}(x)$	
$R(x)$	
$\text{Lock-S}(y)$	
$\text{Unlock}(x)$	
	$\text{Lock-X}(x)$
	$\text{Lock-X}(y)$
	$W(x)$
	$\text{Unlock}(x)$
	$W(y)$
	$\text{Unlock}(y)$
$R(y)$	
$\text{Unlock}(y)$	

T_2 lock y operation
cannot be granted at
this point

T_2 can lock y
after T_1 unlocks y



Why Two Phases?



T_1	T_2
$\text{Lock-S}(x)$ $R(x)$ $\text{Lock-S}(y)$ $\text{Unlock}(x)$	
$R(y)$ $\text{Unlock}(y)$	$\text{Lock-X}(x)$
	$\text{Lock-X}(y)$ $W(x)$ $\text{Unlock}(x)$ $W(y)$ $\text{Unlock}(y)$



Lock Point



- ❖ The transaction's *lock point*
 - ❖ The time point when a transaction has obtained its final lock
- ❖ 2PL schedule is equivalent to a serial schedule where transactions are ordered by their lock points
- ❖ *Question:* Can you find a conflict serializable schedule that cannot be obtained via 2PL?
 - ❖ Yes. (leave to you as an exercise)



Recoverability and 2PL



- ❖ 2PL allows cascading roll-back
 - ❖ 😞 if a transaction is aborted, may need to undo many other transactions
- ❖ How to modify the 2PL protocol to avoid this drawback?
- ❖ Strict two-phase locking
 - ❖ A transaction must hold all its exclusive locks till it commits/aborts
 - ❖ Guarantees cascadeless schedules



Lock Conversion



- ❖ If a transaction T reads an item x first then
- ❖ later writes x , what lock shall T acquire?
 - ❖ [1] T can acquire lock-X on x before it reads x
 - ❖ [2] T acquires lock-S on x before it reads x , then convert lock-S(x) to lock-X(x)
- ❖ The latter case [2] is called *lock conversion*
 - ❖ Allow more concurrency
 - ❖ The DBMS may not know whether a transaction that reads x will or will not write x later



2PL with Lock Conversions



❖ First Phase

- ❖ can acquire a lock-S on item
- ❖ can acquire a lock-X on item
- ❖ can convert a lock-S to a lock-X (upgrade)

❖ Second Phase

- ❖ can release a lock-S
- ❖ can release a lock-X
- ❖ can convert a lock-X to a lock-S (downgrade)

❖ This protocol assures serializability



Lecture Objectives



- ❖ Two-Phase Locking
- ❖ Problems of Locking

- ❖ Timestamp Ordering
- ❖ Optimistic Concurrency Control






Starvation



- ❖ Any locking protocol has to deal with 2 problems:
starvation and *deadlock*
- ❖ Starvation occurs when a lock request will *wait indefinitely*
 - ❖ A transaction is waiting for an X-lock on an item, while many other transactions request and are granted S-lock on the same item.
- ❖ How to prevent starvation?

Deadlock

- ❖ A deadlock occurs when there is a *circular wait* of transactions



- ❖ When a deadlock occurs, the DBMS must abort (and later restart) a transaction in the deadlock
 - ❖ to release all locks held by the aborted transaction



Deadlock



- ❖ Three solutions for solving deadlock
 - ❖ Maintaining the wait-for-graph
 - ❖ A timeout-based scheme
 - ❖ Priority-based scheme

Deadlock Detection

- ❖ A *wait-for graph*

- ❖ A vertex denotes a transaction
- ❖ A directed edge $T_i \rightarrow T_j$ means that T_i is waiting for T_j to release a data item

- ❖ Insert an edge $T_i \rightarrow T_j$ in the graph

- ❖ When T_i requests a data item being held by T_j

- ❖ Remove an edge $T_i \rightarrow T_j$


- ❖ When T_j releases lock of a data item needed by T_i

- ❖ A cycle in the wait-for graph
== deadlock happens

- ❖ Must invoke a deadlock-detection algorithm periodically to detect cycles

T_1	T_2	T_3
Lock-S(x) R(x)		
	Lock-X(y) W(y)	
Lock-S(y)		
		Lock-X(z) W(z) Lock-X(x)
Lock-S(z)		

Wait-for graph:



Note: Don't mix up *precedence graph* and *wait-for-graph*!



Deadlock Recovery



- ❖ How to select a victim (to abort)?
 - ❖ Effort already invested in a transaction
 - ❖ Cost of aborting a transaction (e.g., will it cause cascading aborts?)
 - ❖ Effort to finish the transaction
 - ❖ Age of the transaction
 - ❖ Number of rollbacks (to avoid starvation)
- ❖ Rollback -- determine how far to roll back transaction
 - ❖ More effective to roll back transaction only as far as necessary to break deadlock



Timeout-based Scheme



- ❖ *Timeout*: abort a transaction if it has been waiting too long (for a lock)

- ❖ The timeout period should be:
 - ❖ long enough so that most transactions that are aborted are actually deadlocked
 - ❖ short enough so that deadlocked transactions don't wait too long for their deadlocks to be broken




Deadlock Avoidance Schemes




- ❖ Deadlock avoidance: prevent a transaction to block any transaction with “higher priority”
 - ❖ E.g., assign priorities to transactions based on transaction timestamps (e.g., start time)
 - ❖ If a transaction T_i requests a lock held by a transaction T_j
 - ❖ *Wait-die* scheme:
 - ❖ T_i waits if $T_i.timestamp < T_j.timestamp$
 - ❖ T_i is rolled back otherwise
 - ❖ *Wound-wait* scheme:
 - ❖ T_i waits if $T_i.timestamp > T_j.timestamp$
 - ❖ T_j is rolled back otherwise

Wait-die vs. Wound-wait

- ❖ Wait-die: a young transaction may get rolled back repeatedly



- ❖ Wound-wait: there may be fewer rollbacks



- ❖ With wound-wait, an old transaction never waits.



Cyclic Restart



- ❖ A *cyclic restart* occurs when a transaction is repeatedly rolled back without making progress
- ❖ Both wait-die and wound-wait prevent cyclic restarts because:
 - ❖ The oldest transaction never gets rolled back
 - ❖ The oldest transaction eventually gets to execute to completion
 - ❖ A transaction will eventually become the oldest transaction in the system
- ❖ *Question:* when a transaction is restarted, shall we use its old timestamp, or get a new timestamp?



Summary of These Schemes



- ❖ The wait-for-graph (WFG)
 - ❖ its maintenance requires too much resource, especially if deadlock occurs rarely
- ❖ A timeout-based scheme
 - ❖ easy to implement
 - ❖ E.g., if an application does not receive any terminal input before the timeout, it may roll back a transaction
- ❖ Priority-based schemes
 - ❖ may cause many unnecessary rollbacks



Lecture Objectives



- ❖ Two-Phase Locking
- ❖ Problems of Locking



- ❖ Timestamp Ordering
- ❖ Optimistic Concurrency Control





Timestamp Ordering



- ❖ Each transaction T is given a *timestamp* $\text{TS}(T)$
 - ❖ An old transaction (with small timestamp) has high priority
- ❖ Given two transactions T_i and T_j , if $\text{TS}(T_i) < \text{TS}(T_j)$, we want a schedule S such that S is equivalent to a serial schedule S' with T_i before T_j
- ❖ Maintain 2 timestamps for each data item Q :
 - ❖ $W\text{-}\text{TS}(Q)$: the largest timestamp of any transaction that has executed $\text{write}(Q)$
 - ❖ $R\text{-}\text{TS}(Q)$: the largest timestamp of any transaction that has executed $\text{read}(Q)$

- ❖ TOP Idea: ensure that conflicting operations are executed in timestamp order
 - ❖ Note: a restarted transaction will be given a *new timestamp*
- ❖ On a transaction T issuing $\text{read}(Q)$:
 - ❖ $\text{TS}(T) < \text{W-TS}(Q)$: T issues $\text{read}(Q)$ too late \Rightarrow restart T
 - ❖ $\text{TS}(T) > \text{W-TS}(Q)$: execute $\text{read}(Q)$ and update $\text{R-TS}(Q)$ to $\max\{\text{R-TS}(Q), \text{TS}(T)\}$

time	T_1	T_2
1	<i>Start</i>	
2		<i>Start</i>
3		$\text{Write}(Q)$
4	$\text{Read}(Q)$	

$$\text{TS}(T_1)=1 \quad \text{W-TS}(Q)=2$$

time	T_1	T_2
1		<i>Start</i>
2	<i>Start</i>	
3		
4		$\text{Read}(Q)$

$$\text{TS}(T_1)=2 \quad \text{W-TS}(Q)=1 \quad 31$$

❖ On a transaction T issuing $\text{write}(Q)$:

- ❖ If $\text{TS}(T) < \text{R-TS}(Q)$ or $\text{TS}(T) < \text{W-TS}(Q)$
 T issues $\text{write}(Q)$ too late \Rightarrow restart T
- ❖ Else: execute $\text{write}(Q)$,
update $\text{W-TS}(Q)$ to $\max\{\text{W-TS}(Q), \text{TS}(T)\}$

time	T_1	T_2
1		<i>Start</i>
2	<i>Start</i>	
3	$\text{Read}(Q)$	
4		$\text{Write}(Q)$

$$\text{TS}(T_2)=1 \quad \text{R-TS}(Q)=2$$

time	T_1	T_2
1		<i>Start</i>
2		<i>Start</i>
3		$\text{Write}(Q)$
4		$\text{Write}(Q)$

$$\text{TS}(T_2)=1 \quad \text{W-TS}(Q)=2$$



TOP Properties



- ❖ Ensure conflict serializability 😊
 - ❖ We skip the proof

- ❖ Deadlock free 😊
- ❖ May lead to starvation of long transactions 😞
- ❖ Schedules can be un-recoverable 😞




TOP: Managing Timestamps



- ❖ Large space requirement 😞
 - ❖ need 2 timestamps for each data item
- ❖ Be careful when we manipulate timestamps
 - For example:
 - ❖ Store timestamps on non-volatile storage or not?
 - ❖ When a transaction aborts, do we rollback a timestamp's value or not?
 - ❖ A read operation causes a write to a data item's timestamp

- ❖ In TOP, when a transaction T issues a $\text{write}(Q)$, if $\text{TS}(T) < \text{W-TS}(Q)$, then restart T
- ❖ However, restarting T is really unnecessary
- ❖ If no transaction should have read T' 's value,
then ignore T' 's write

How to check this condition?



- ❖ TOP with *Thomas write rule* may produce schedules that are view serializable but not conflict serializable
- ❖ Example:

Time	T_1	T_2	
1	<i>Start</i>		$TS(T_1) = 1$
2		<i>Start</i>	$TS(T_2) = 2$
3	Read(Q)		
4		Write(Q)	$R-TS(Q) = 1$
5		Write(Q)	$W-TS(Q) = 2$



Lecture Objectives



- ❖ Two-Phase Locking
- ❖ Problems of Locking

- ❖ Timestamp Ordering
- ➡ ❖ Optimistic Concurrency Control





Optimistic Concurrency Control (OCC)



- ❖ Both 2PL and TOP have bookkeeping overhead for each read/write operation.
 - ❖ Wasteful if conflicts occur rarely in the system
- ❖ Also, lock-based protocols resolve conflicts by blocking
 - ❖ Could lead to severe *data contention thrashing*:
 - adding more transactions to the system
 - >more transactions are waiting for locks
 - >transactions thus spend time waiting rather than working
 - >decreases the system's throughput



Optimistic Concurrency Control



- ❖ Optimistic concurrency control *assumes that conflicts don't happen frequently*
 - ❖ Allow transactions to proceed first
 - ❖ When a transaction commits, the OCC scheduler verifies that the transaction is not in any conflict
- ❖ OCC would be efficient if most of the transactions are read-only




The OCC Protocol




- ❖ OCC requires each transaction to declare its *ReadSet* and *WriteSet* before execution
- ❖ A transaction T is divided into 3 phases.
 - ❖ *Read phase*:
 - ❖ T reads into local memory the data items it reads/writes
 - ❖ Writes to these data items are done only to the local memory
 - ❖ *Validation phase*:
 - ❖ When T is ready to commit, the scheduler checks if T conflicts with any other transactions
 - ❖ If so, the scheduler rejects T's commit and restarts T
 - ❖ *Write phase*:
 - ❖ T passes its validation test
 - ❖ The values T wrote to the local copies are copied to the database


- ❖ Each transaction T is associated with 3 timestamps:
 - ❖ $\text{Start}(T)$: start time of T
 - ❖ $\text{Validation}(T)$: time at which T starts its validation
 - ❖ $\text{Finish}(T)$: time at which T completes its write phase
- ❖ *Goal of the validation test*: order transactions with conflicting operations based on their validation times



- ❖ Consider two transactions T and T' such that $\text{validation}(T') < \text{validation}(T)$.
- ❖ OCC would order conflicting operations of T and T' such that those of T' precede those of T .
- ❖ Consider validating T :
 - ❖ Case 1: $\text{Finish}(T') < \text{Start}(T)$
 - ❖ All conflicting operations (if exist) are ordered properly



- ❖ Case 2: $\text{Start}(T) < \text{Finish}(T') < \text{Validation}(T)$
 - ❖ Since all reads by T' occur before $V(T')$ and all writes by T occur after $V(T)$ and $V(T') < V(T)$, any *read-write conflicts* are ordered properly.
 - ❖ Since all writes by T' occur before $F(T')$ and all writes by T occur after $V(T)$ and $F(T') < V(T)$, any *write-write conflicts* are ordered properly




— reads
— writes

Let's use *x-y conflict* to denote an x operation executed by T' and a y operation executed by T


- ❖ Case 2: $\text{Start}(T) < \text{Finish}(T') < \text{Validation}(T)$
 - ❖ However, *write-read conflicts* may violate the (intended) serialization order.
 - ❖ Solution: if $\text{ReadSet}(T) \cap \text{WriteSet}(T') \neq \emptyset$, abort T.

A write by T' could follow a read by T !




❖ Case 3: $\text{Validation}(T) < \text{Finish}(T')$

- ❖ Technically, the system does not know $F(T')$ yet, since it is validating T . The system, however, knows that $F(T') > V(T)$ because T' has not finished its write phase.
- ❖ Since all reads by T' occur before $V(T')$ and all writes by T occur after $V(T)$ and $V(T') < V(T)$, any *read-write conflicts* are ordered properly.



- ❖ Case 3: $\text{Validation}(T) < \text{Finish}(T')$
 - ❖ *Write-write conflicts* could violate the (intended) serialization order.
 - ❖ Solution: if $\text{WriteSet}(T) \cap \text{WriteSet}(T') \neq \emptyset$, abort T.
 - ❖ *Write-read conflicts* could violate the (intended) serialization order.
 - ❖ Solution: if $\text{ReadSet}(T) \cap \text{WriteSet}(T') \neq \emptyset$, abort T.





OCC: Validation Rules



- ❖ If $\text{Finish}(T') < \text{Start}(T)$
 - ❖ T passes its validation
- ❖ If $\text{Finish}(T') > \text{Start}(T)$
 - ❖ Write-read conflicts could violate the serialization order
 - ❖ Write-write conflicts could violate the serialization order if $\text{Finish}(T') > \text{Validation}(T)$
- ❖ To validate a transaction T , **check** these two rules for every transaction T' with $\text{Validation}(T') < \text{Validation}(T)$
 - ❖ Rule 1: if $\text{Finish}(T') > \text{Start}(T)$ then
if $\text{ReadSet}(T) \cap \text{WriteSet}(T') \neq \emptyset$, abort T
 - ❖ Rule 2: if $\text{Finish}(T') > \text{Validation}(T)$ then
if $\text{WriteSet}(T) \cap \text{WriteSet}(T') \neq \emptyset$, abort T



Conclusions



Two classes of concurrency control protocols:

- ❖ Pessimistic (e.g., 2PL and TOP)
- ❖ Optimistic (e.g., OCC)
 - ❖ performs well for read-intensive transactions
 - ❖ but validation and restart are more costly



谢谢！

DBGroup @ SUSTech

Dr. Bo Tang (唐博)

tangb3@sustech.edu.cn

