



2025 Fall CSE5025

Combinatorial Optimization

组合优化

Instructor: 刘晟材

Lecture 5-2: Integer Programming for Combinatorial Optimization – Part II

Agenda for Today's Lecture

In this lecture, we will focus on

- The core idea of Branch and Cut (B&C)
- Constructive Heuristics
- Local Search
- Metaheuristics

Learning objectives for this lecture

- Know the performance bottleneck of B&B
- Know the basic idea of B&C
- Master the differences between different types of heuristics

Notations:

- Z_{IP} : the optimal value of the original IP (the one we want to solve).
- Z_{LP} : the optimal value of the LP relaxation of the original IP.
- Z_{INC} : the value of the incumbent solution (the best-so-far integer solution)
- Z_{LP}^P : the optimal value of the LP relaxation at node P in the B&B tree

We always have (for maximization problem):

- $Z_{INC} \leq Z_{IP} \leq Z_{LP}$
- $Z_{LP}^P \leq Z_{LP}$

Recap: The B&B Algorithm

In our last lecture, we saw that Branch & Bound is an intelligent search. Its efficiency crucially depends on one thing: **Pruning** (剪枝).

- We maintain an incumbent (the best-so-far integer solution) Z_{INC} .
- At each node P , we solve the LP relaxation to get its optimal value Z_{LP}^P , which is an optimistic estimate of any integer solution than can be **obtained in this entire branch (any subproblem created from node P)**.
- The pruning (by bound) rule (for maximization problem):

$$\text{IF } Z_{LP}^P \leq Z_{INC} \text{ THEN Prune}$$

Why is B&B Sometimes Inefficient?

The Problem: A “Weak” or “Loose” Bound. The pruning rule works if Z_{LP} is close to the true integer optimum Z_{IP} .

The Integrality Gap (IG): The gap between the LP relaxation and the true integer solution at the root node:

- (for minimization): $Z_{IP} - Z_{LP}$
- (for maximization): $Z_{LP} - Z_{IP}$
- the ratio definition Z_{LP}/Z_{IP} is also used in the literature

A large IG means the initial bound is very loose.

The Consequence of a Weak Bound

The core issue is that Z_{LP}^P must drop low enough (for maximization problem) to compete with the Incumbent Z_{INC} .

The Bound Monotonicity: As the B&B tree goes deeper, Z_{LP}^P is non-increasing because we are only adding more constraints (branching).

$$Z_{LP}^P \leq Z_{LP}^{parent\ node\ of\ P}$$

The Challenge of a Loose Root:

- Suppose $Z_{IP} = 100$ and $Z_{INC} = 95$
- If the root bound $Z_{LP} = 1000$ (a loose bound), the entire tree starts high.
- For the pruning rule $Z_{LP}^P \leq Z_{INC} = 95$ to activate, the B&B algorithm has to spend enormous resources (branching) to bring Z_{LP}^P down from the “ceiling” (1000) to a range (≤ 95) where effective pruning can finally occur. Without a tight initial bound, the B&B engine is severely hampered.

The Core Idea: Cutting Planes



A cutting plane (割平面), or “cut”, is a special constraint that we **add to our LP**.


A “cut” is a **valid inequality** with two properties:

- It is **satisfied by all feasible integer solutions**. That is to say, it never cuts off the true integer optimum Z_{IP} .
- It is **violated by the current fractional solution x^*** . In other words, it always cuts off the current “bad” optimal solution of the LP).


The new LP will have **a smaller feasible region** than the previous LP and hence has **a tighter bound**.

The process of finding and adding these cuts is called “**separation**”.


Visualizing a Cutting Plane




(1)



(2)



(3)



(4)

B&C improves the bound at each node, which in turn improves pruning.

Weak Bound (Original B&B):

- In this scenario, the root node Z_{LP} was very loose, and the bound at node P is still too loose.
- $Z_{LP}^P = 10 \dots Z_{IP} = 100 \dots Z_{INC} = 110$
- $Z_{LP}^P \not\geq Z_{INC} \rightarrow$ No Pruning.

An example (minimization problem)

Tightened Bound (Branch & Cut):

- We add cuts...
- $Z_{LP}^{P,I} = 50 \dots Z_{IP} = 100 \dots Z_{INC} = 110$
- $Z_{LP}^{P,I} \not\geq Z_{INC} \rightarrow$ No Pruning.
- We add more cuts...
- $Z_{LP}^{P,II} = 105 \dots Z_{IP} = 100 \dots Z_{INC} = 110$
- $Z_{LP}^{P,II} \not\geq Z_{INC} \rightarrow$ No Pruning.
- We add even more cuts...
- $Z_{LP}^{P,III} = 115 \dots Z_{IP} = 100 \dots Z_{INC} = 110$
- $Z_{LP}^{P,III} \geq Z_{INC} \rightarrow \text{PRUNE!}$

Algorithmic Steps of B&C

B&C integrates the “cut loop” inside the B&B algorithm.

- 1) Initialize a “global incumbent” $Z_{INC} = -\infty$.
- 2) Add the root LP relaxation to a “**set**” of problems.
- 3) While **set** is not empty:
 - a. Select a problem P from the set.
 - b. Solve its LP relaxation, get Z_{LP}^P and \boldsymbol{x}^* .
 - c. Pruning 1 (by Bound): If $Z_{LP}^P \leq Z_{INC}$, prune this branch, and goto step 3).
 - d. Pruning 2 (Infeasible): If $Z_{LP}^P = -\infty$, prune this branch and goto step 3).
 - e. If \boldsymbol{x}^* are all integers, we have found a feasible integer solution. Check if $Z_{LP}^P > Z_{INC}$ and update the incumbent Z_{INC} accordingly. Goto Step 3).

Algorithmic Steps of B&C

f. **Cut (the new part):** If x^* is fractional.

- Enter the “Cut Loop”:
 - **Find Cuts:** Run “separation algorithms” to find valid cuts that x^* violates.
 - **If Cuts Found:** Add them to the LP and go back to Step 3-b (Solve).
 - **If No Cuts Found:** Exit the Cut Loop.

e. **Branch:** If we exit the Cut loop and x^* is still fractional, then we create two new subproblems with added constraints, respectively. Add both to the set. Goto Step 3).

4) Return Z_{INC} .

B&C Applied to the DFJ formulation of TSP

The DFJ Formulation (Dantzig-Fulkerson-Johnson) is a strong IP model for TSP, providing a very tight LP relaxation bound.

- Model Strength: The LP Relaxation (Z_{LP}) of the full DFJ formulation is extremely tight, meaning Z_{LP} is very close to Z_{IP} . This provides excellent pruning potential.
- The Constraint Issue: DFJ achieves this strength by including Subtour Elimination Constraints (SECs). There are $O(2^n)$ such constraints:

$$\sum_{i \in S} \sum_{j \notin S} x_{ij} \geq 1 \text{ for all } S \subseteq V, S \neq \emptyset, S \neq V$$

The Paradox: The strongest model is **computationally difficult** (impossible) to load into the solver upfront.

The Full DFJ Formulation of TSP

$$\text{Minimize } Z = \sum_{i=1}^n \sum_{j \neq i} c_{ij} x_{ij}$$

$$\text{s. t. } \sum_{i \neq j} x_{ij} = 1 \text{ for each } j = 1, \dots, n$$

$$\sum_{j \neq i} x_{ij} = 1 \text{ for each } i = 1, \dots, n$$

$$\sum_{i \in S} \sum_{j \notin S} x_{ij} \geq 1 \text{ for all } S \subseteq V, S \neq \emptyset, S \neq V$$

$$x_{ij} \in \{0,1\} \text{ for all } i \neq j$$

The Core Idea: Lazy Constraint Generation

B&C resolves this paradox by treating the $O(2^n)$ constraints as “Lazy Constraints” that are generated on demand.

Algorithmic steps at a node:

- 1) **Start:** The LP contains the basic assignment constraints (enter/leave once) and only those SECs inherited from its parent, and a branching constraint.
- 2) **Solve LP:** The solver returns a fractional solution x^* (which contains subtours).
- 3) **Separate (Cut Loop):** We run an efficient **separation algorithm** (e.g., min-cut) on x^* to find a violated SEC (a subtour S), e.g., for a subtour $1 \rightarrow 2 \rightarrow 1$, we identify $S = \{1, 2\}$.
- 4) **Add Cut:** We add only the one violated constraint for that S back to the LP model.
- 5) **Re-Solve & Repeat:** The solver is forced to find a new, valid fractional solution x^{**} that avoids S . This loop continues until no more violated SECs can be found.

Pure B&B is slow if LP relaxation is “weak”. B&C = B&B + Cutting Planes.

How does it work?

- At each node P , it solves the LP, finds violated cuts, adds those cuts, and re-solves. This tightens the bound Z_{LP}^P

Why is it efficient?

- A tighter bound \rightarrow more pruning.
- It allows us to use “strong” but exponentially large formulations (like TSP’s DFJ) in a computationally feasible way. Note for TSP, more powerful cuts than the SECs (cuts) are used.

B&C is the core technology inside modern IP solvers (Gurobi, CPLEX) and specialized solvers (Concorde for TSP).

Lecture 6-1: Heuristics and Metaheuristics for Combinatorial Optimization

Recap: We just learned B&C. It can solve TSP optimally.

The Reality:

- B&C works for TSP up to thousands, but it takes massive computing power.
- For many other NP-hard problems, exact methods fail at $n = 50$ or $n = 100$.

Real-World Needs:

- Scale: What if we have 1000000 cities?
- Time: What if we need an answer in seconds, not days?
- Complexity: What if the constraints are messy and hard to model in IP?

The **Trade-off**: We sacrifice Optimality for Speed. We seek a “Good Enough” solution in “Reasonable Time”.

Approximation Algorithm (Previous Lectures):

- A polynomial-time algorithm with a provable performance guarantee. e.g., Christofides for Metric TSP: always within $1.5 \times OPT$.
- **Focus: Worst-case analysis.**

Heuristic (This Lecture):

- An algorithm designed to find good solutions empirically.
- **No provable guarantee** on solution quality (in the general case).
- It might find the optimal solution, or it might be far off.
- **Focus: Average-case empirical performance and practical speed.**

Metaheuristic (This lecture):

- A general-purpose high-level framework or strategy to guide heuristic search.

Constructive Heuristics:

- Start from an empty solution.
- Build the solution piece by piece (greedily).
- **Pros:** Extremely fast, easy to implement.
- **Cons:** Often lower quality, decisions are myopic.

Improvement Heuristics (Local Search):

- Start from a complete solution.
- Iteratively make small changes to improve it.
- **Pros:** Better quality than constructive methods.
- **Cons:** Susceptible to getting trapped in Local Optima.

Nearest Neighbor (NN):

- Logic: Always travel to the closest unvisited city.
- Weakness: Forced to take very long edges at the end to close the tour.

Cheapest Insertion (CI):

- Logic: Find the unvisited city k and the specific position (i, j) in the current tour that minimizes the increase in tour cost.
- Strength: Directly optimizes the objective function at each step.

Farthest Insertion (FI):

- Logic: Start with a subtour. Insert the node that is farthest from the current tour to minimize cost increase.
- Strength: Establishes the general shape of the tour early

Constructive Example: Graph Coloring

Problem: Assign a color to each vertex of a graph such that no adjacent vertices share the same color, minimizing the total number of colors.

The “Welsh-Powell” Heuristic:

- Calculate Degrees: Compute the degree of every vertex.
- Sort: Order vertices by degree in descending order.
 - Rationale: High-degree nodes are the hardest to color.
 - Coloring Loop:
 - Assign the first available color (Color 1) to the head of the list.
 - Traverse the list: Color any non-adjacent uncolored node with Color 1.
 - Move to Color 2, repeat until all nodes are colored.

Improvement Heuristics: Local Search

Concept: Instead of building a solution from scratch, we improve or repair an existing one.

The General Algorithm:

1. **Initialization:** Generate an initial solution S (random or constructive).
2. **Evaluation:** Calculate the cost objective $f(S)$.
3. **Neighborhood Search:** Examine solutions S' that are "near" S .
4. **Move:** If a neighbor S' exists such that $f(S') < f(S)$:
 - Update current solution: $S \leftarrow S'$.
 - Repeat step 3.
5. **Termination:** If no neighbor is better, STOP.
 - The current solution S is a **Local Optimum**.

Defining the Neighborhood Structure

The definition of ‘‘Neighborhood’’ $N(S)$ defines the search landscape. $N(S)$ is the set of solutions reachable from S by a move operator.

Common Neighborhoods for TSP:

- 2-Opt: Delete 2 edges and reconnect the path (reverse a tour segment). Size: $O(n^2)$.
- Swap: Exchange the positions of two cities. Size: $O(n^2)$.
- Relocate (Insert): Move a city to a new position in the tour. Size: $O(n^2)$.

Common Neighborhoods for Binary Problems (e.g., Knapsack):

- Bit Flip: Change x_i from 0 to 1 (or 1 to 0). Size: $O(n)$.
- Pair Swap: Exchange a selected item with an unselected one. Size: $O(n^2)$.

Illustration of 2-OPT








Illustration of 3-OPT




Original tour : ABC




Method 1. ACB




Method 2. ACB⁻¹




Method 3. AC⁻¹B




Method 4. AC⁻¹B⁻¹



Method 5. ABC⁻¹



Method 6. AB⁻¹C



Method 7. AB⁻¹C⁻¹

Search Strategy 1: First Improvement

Definition: Scan neighbors in $N(S)$ sequentially or randomly and accept the first one found that improves the objective function.

Algorithm:

1. Order the neighbors in $N(S)$ (e.g., random permutation).
2. Evaluate neighbor S' .
3. If $f(S') < f(S)$, move to S' immediately.
4. Restart search from the new solution.

Pros & Cons:

- **Pro:** Very fast per iteration. Ideal for large neighborhoods.
- **Con:** The improvement gain at each step might be small.

Search Strategy 2: Best Improvement

Definition: Evaluate **ALL** neighbors in $N(S)$ and move to the one that yields the maximum improvement in objective function (e.g., reduction in cost).

Pros & Cons:

- **Pro:** Maximizes the gain at every single step; Deterministic path for a given starting point
- **Con:** Very slow per iteration if $|N(S)|$ is large (e.g., $O(n^3)$).

Algorithm:

1. Set `best_neighbor` = null, `min_cost` = $f(S)$ (Current Cost).
2. **For every** neighbor S' in $N(S)$:
 - If $f(S') < min_cost$:
 - Update `min_cost` = $f(S')$.
 - Update `best_neighbor` = S' .
3. **Move:**
 - If `best_neighbor` is not null (Improvement found):
 - Update $S \leftarrow best_neighbor$.
 - **Else** (No neighbor is better):
 - **STOP.** Current S is a Local Optimum.

Comparison of Two Improvements

The choice determines the **trade-off between speed per step and gain per step.**

| Feature | First Improvement (First Accept) | Best Improvement (Steepest Descent) |
|---------------------------|---|--|
| Move | Moves to the first neighbor S' found | Moves to the neighbor S' that yields the |
| Acceptance | such that $f(S') < f(S)$. | maximum improvement in the entire neighborhood. |
| Iteration Behavior | Stops scanning the current neighborhood immediately once an improving move is found. | Must scan the entire neighborhood $N(S)$ completely before making a move. |
| Computational Cost | Low cost per iteration (often just a few checks). | High cost per iteration (requires \$ |
| Convergence Path | Takes many small steps. Path depends on check order (stochastic). | Takes fewer, larger steps. Path is deterministic. |

When to Choose Which Strategy

The optimal choice depends on the neighborhood size and time budget.

- Prioritize first improvement for large neighborhoods
 - If your move operator (e.g., 3-Opt for TSP) generates $O(n^3)$ or more neighbors, running best improvement becomes prohibitively slow.
 - Randomization is important: To prevent the search from following a predictable, low-quality path, the order of checking neighbors in first improvement should be randomized at the start of each new iteration.
- Use best improvement for solution refinement on small neighborhoods.
- The Time-to-Quality Curve (Empirical):
 - While best improvement takes fewer steps, first improvement often achieves a high-quality solution faster in terms of wall-clock time.

Impact of Neighborhood Size

Small Neighborhoods (e.g., 2-Opt, 1-Flip):

- Size: Polynomial, usually low degree ($O(n)$, $O(n^2)$)
- Speed: Fast to evaluate each step.
- Landscape: “Rugged” with many small valleys.
- Result: The search gets stuck in local optima very easily.

Large Neighborhoods (e.g., 3-Opt):

- Size: Larger polynomial $O(n^3)$ or exponential.
- Speed: Slow to evaluate.
- Landscape: “Smoothen”. The move allows jumping over small hills.
- Result: Finds higher quality local optima but is computationally expensive.

Implementation Tip: Delta Evaluation



Naive Evaluation:

- Recalculate the objective function $f(S)$ from scratch after every move.
- Cost: $O(n)$ per neighbor.

Delta (Δ) Evaluation:

- Calculate only the change caused by the move.
- Example (TSP 2-Opt):
 - Remove edges (A,B) and (C,D). Add (A,C) and (B,D).
 - $\Delta = dist(A,C) + dist(B,D) - dist(A,B) - dist(C,D)$.
- Cost: $O(1)$ per neighbor.


Impact: This speeds up Local Search by a factor of n . This optimization is critical for solving large problem instances.

The Fundamental Problem: Local Optima

Definition: A solution S^* is a Local Optimum if $f(S^*) \leq f(S')$ for all neighbors $S' \in N(S^*)$.

The Trap:

- Local Search algorithms are Greedy.
- They only accept improving moves.
- Once they hit a local optimum (a “valley”), they stop.
- They cannot temporarily climb “uphill” to reach a deeper valley (global optimum) located elsewhere in the search space.
- To solve this, we need **Metaheuristics**.



Definition: A high-level strategy that guides other heuristics to search for solutions beyond local optimality.

Key Feature: They allow **non-improving moves** (deterioration) to escape local optima.

The Core Duality:

- Exploitation (Intensification):
 - Searching carefully around good solutions to refine them.
 - Mechanism: local search (Descent).
- Exploration (Diversification):
 - Visiting entirely new regions of the search space.
 - Mechanism: Random moves, large jumps, restarts.

Trajectory-Based (a single solution)

- The algorithm maintains and modifies a **single solution** at a time.
- It traces a “trajectory” or path through the search space.
- **Focus: Depth-oriented. Good at Exploitation.**
- **Examples: Simulated Annealing (模拟退火), Tabu Search (禁忌搜索).**

Population-Based (A set/population of solutions)

- The algorithm maintains a set (population) of solutions.
- Solutions interact, combine, or compete to create new solutions.
- **Focus: Breadth-oriented. Good at Exploration.**
- **Examples: Genetic Algorithms (遗传算法).**

Trajectory Method: Simulated Annealing (SA)



Inspiration: Physical annealing in metallurgy. Heating metal and cooling it slowly allows atoms to settle into a strong, low-energy crystal structure.

Analogy:

- **Energy**: Objective Function Cost $f(S)$.
- **Temperature**: A control parameter determining the probability of accepting bad moves.

Core Idea:

- To escape local optima, we must occasionally accept uphill (worse) moves.
- The probability of accepting a worse move should depend on the size of the deterioration and the current “Temperature”.

Simulated Annealing: The Algorithm

1. Initialize solution S and Temperature T .
2. **While** $T > T_{min}$:
3. **Repeat** L times (Inner Loop):
 - Pick a random neighbor $S' \in N(S)$.
 - Calculate $\Delta = f(S') - f(S)$.
 - **If** $\Delta < 0$ (**Improvement**):
 - Always accept: $S \leftarrow S'$.
 - **If** $\Delta \geq 0$ (**Worsening**):
 - Accept with probability $P = e^{-\Delta/T}$.
 - (Generate random $r \in [0, 1]$. If $r < P$, move to S').
4. **Cool Down:** Decrease T (e.g., $T \leftarrow 0.99 \cdot T$).

SA: The Logic of temperature

The Metropolis Criterion: $P(\text{accept}) = e^{-\Delta/T}$

High Temperature ($T \rightarrow \infty$):

- $e^{-\Delta/T}$ is close to 1.
- The algorithm accepts almost all moves, even terrible ones.
- Behavior: random walk (pure exploration).

Low Temperature ($T \rightarrow 0$)

- $e^{-\Delta/T}$ is close to 0.
- The algorithm rejects almost all worsening moves.
- Behavior: local search / hill climbing (pure exploitation).

Strategy: Start with high exploration to find the right region, then transition to high exploitation to refine the solution.

Philosophy: “Intelligent” search using **Memory** rather than randomness.

Core Mechanism:

- 1) Generate all (or many) neighbors of the current solution S .
- 2) Move to the best neighbor S' , even if it is worse than S .
 - This automatically allows climbing out of local optima.

The Cycling Problem:

- If we move $S \rightarrow S'$ (worse), the best neighbor of S' is usually S (the previous optimum).
- The search will cycle $S \rightarrow S' \rightarrow S$ indefinitely
- Solution: We need a **Tabu List**.

Short-Term Memory (The Tabu List):

- Records recent moves or solution attributes.
- **Rule:** A move is **Tabu** (forbidden) if it is in the list.
- Effect: Forces the search to be away from the current local optimum (prevents cycling).

Aspiration Criteria (渴望准则、特赦准则):

- A rule to override the Tabu status.
- Rule: If a move is **Tabu**, but it produces a solution better than the global best found so far, allow it anyway.

Long-Term Memory:

- Frequency-based: Track how often edges/nodes are used.
- Penalize frequently used features to force the search into unexplored regions.

Inspiration: Darwinian Natural Selection.

Key Difference from single solution-based methods: operates on a **population of solutions**, not just one.

Terminology:

- Chromosome (染色体): Representation of a solution (e.g., permutation [1, 4, 2, 3]).
- Fitness (适应度): The objective function value (quality of the solution).
- Generation (代): One iteration of the algorithm.

Hypothesis: By combining the “genetic material” (sub-structures) of two high-quality parent solutions, we may create an offspring solution that inherits the strengths of both.

1. **Initialization:** Generate a random population of N individuals.
2. **Evaluation:** Compute Fitness for every individual.
3. **Selection:** Choose "parents" for the next generation.
 - Biased towards better individuals (Survival of the Fittest).
4. **Crossover (Recombination):** Combine Parent A and Parent B to create a Child.
 - This is the **primary** search operator.
5. **Mutation:** Randomly tweak the child (e.g., flip a bit, swap cities).
 - Maintains diversity and prevents premature convergence.
6. **Replacement:** Select survivors to form the next generation.

Goal: Select parents such that better individuals have a higher chance, but weak ones are not totally excluded (to maintain diversity).

Roulette Wheel Selection (轮盘赌选择):

- Probability of selecting S_i is proportional to its fitness: $P(S_i) = f(S_i)/\sum f(S_j)$
- **Risk:** One super-individual can dominate early, reducing diversity.

Tournament Selection (锦标赛选择):

- Randomly pick k individuals from the population.
- Select the best one among the k to be a parent.
- Tunable: Large k = High pressure (Exploitation). Small k = Low pressure (Diversity or Exploration).

Binary Crossover (e.g., Knapsack):

- Parent A: 11111 | 00000
- Parent B: 00000 | 11111
- **One-Point Crossover:** Cut at a random line, swap tails.
- Child: 1111111111 (Inherits first half of A, second half of B).

Permutation Crossover (e.g., TSP):

- **Challenge:** Standard crossover creates invalid tours (duplicates/missing cities).
- **Ordered Crossover (OX1):**
 - Copy a random segment from Parent A to the Child.
 - Fill remaining slots with cities from Parent B; **crucially, fill them in the order they appear in Parent B, skipping cities already in the segment.**
- This preserves relative ordering information.

Mutation (变异):

- Role: diversity maintenance.
- Without mutation, GA can only shuffle existing genes. If a necessary gene is missing from the entire population, crossover cannot recover it.
- Mechanism: Randomly modify the child with low probability.
- TSP Example: Randomly swap two cities in the tour.

Elitism (精英):

- Problem: Crossover and Mutation are destructive. We might accidentally destroy the best solution found so far.
- Mechanism: Always copy the top $k\%$ of individuals from the current generation directly to the next generation, unchanged.
- Guarantees that the best solution quality never decreases.

The complementary strengths:

- GA: Excellent at global exploration. They find the “right hills” but are slow at climbing to the exact peak.
- Local Search: Excellent at local exploitation. It climbs peaks very fast but gets stuck on the wrong hill.

Memetic Algorithm (GA + Local Search):

- Strategy: Apply local search to some solutions generated by crossover/mutation:
 - 1) Child = Crossover(Parent A, Parent B).
 - 2) Child = Mutation(Child).
 - 3) Child = Local_Search(Child).

Memetic algorithms are the state-of-the-art for many combinatorial problems.

Constructive Heuristics:

- Use for initialization or extremely time-constrained applications.

Local Search:

- Use to refine constructive/random solutions.
- Implement with Delta Evaluation for speed.

Metaheuristics:

- Simulated Annealing: best for problems with messy constraints. Easy to implement.
- Tabu Search: excellent for routing (TSP, VRP) and scheduling.
- Genetic Algorithms and Memetic Algorithms: The typical high-performance choice for hard problems, combining evolutionary diversity with local optimization.