

ToT-Prompt Optimization for LLM-based Automatic Summarization

Reporter: Zhao Zhao Xiaoqi Qiu
Date: 2025/12/10

CATALOGUE

CATALOGUE

01

Introduction

A brief introduction
to this problem

02

Formulization

- Math Modeling
- Complexity

03

Method Design

- Design Rationals
- Method

04

Experiments

- Dataset
- Results

01

Introduction

Introduction

📌 Core Problem

Select and arrange prompt words, maximize the quality of summaries generated by LLMs

🔍 Problem Nature

- Combinatorial Optimization Problem
- Huge Search Space: n components $\rightarrow O(n!)$ possible permutations and combinations
- High Evaluation Cost

Introduction



Practical Application Background

- Financial Analysis:
Efficiently process financial news, corporate reports
- Scientific Research:
Accelerating knowledge extraction and integration
- Law & Healthcare:
Quickly condense legal cases, generate medical record summaries

02

Formulization

Formalization of the Problem

Find an optimal ordered sequence from n prompt components within a limited evaluation budget.

Maximize: $Q(\sigma)$

where $Q(\sigma) = f_{comp}(\sigma) + f_{order}(\sigma) - p(\sigma)$ Component Base Score, Order Synergy Effect, Penalty Terms

Decision variables: $\sigma = (\sigma_1, \sigma_2, \dots, \sigma_k) \in \Pi_k(C)$

Let $C = c_1, c_2, \dots, c_n$ be the set of all available prompt components.

The decision variable is an ordered sequence:

$$\sigma = (\sigma_1, \sigma_2, \dots, \sigma_k)$$

where:

$k \leq n$: sequence length (variable)

$\sigma_i \in C$: component ID at position i in the sequence

$\Pi_k(C)$: set of all permutations of k elements from C

Constraints

1. Dependency Constraints

For each component $c \in C$, let its dependency set be $D(c) \subseteq C$:

$$\forall c \in \sigma, \forall d \in D(c) : d \in \sigma \text{ and } pos(d) < pos(c)$$

where $pos(x)$ represents the position index of component x in sequence σ .

2. Mutual Exclusion Constraints

For mutually exclusive component pairs $(c, d) \in X$ (where X is the mutual exclusion relation set):

$$c \notin \sigma \text{ or } d \notin \sigma$$

Constraints

3. Length Constraints

Let $t(c)$ be the token count of component c , and $t_{\text{conn}}(c_i, c_j)$ be the connection token count between components:

$$\sum_{i=1}^k t(\sigma_i) + \sum_{i=1}^{k-1} t_{\text{conn}}(\sigma_i, \sigma_{i+1}) \leq T_{\max}$$

4. Position Constraints

For components with position requirements:

Start components: $c \in C_{\text{start}} \Rightarrow \text{pos}(c) = 1$

End components: $c \in C_{\text{end}} \Rightarrow \text{pos}(c) = k$

Middle components: $c \in C_{\text{middle}} \Rightarrow 1 < \text{pos}(c) < k$

NP-Hard Proof

Theorem: The prompt sequence optimization problem is NP-hard.

Proof by reduction from Directed Hamiltonian Path (DHP):

Given: A directed graph $G = (V, E)$ with $|V| = n$ vertices.

Construct an instance of our problem:

Each vertex $v_i \in V$ corresponds to a prompt component c_i

For each edge $(v_i, v_j) \in E$, set synergy $s(c_i, c_j) = 1$

For non-edges $(v_i, v_j) \notin E$, set $s(c_i, c_j) = -\infty$

Set all component base scores $w(c_i) = 0$

Set length constraint to accept all sequences

Set no dependency or mutual exclusion constraints

NP-Hard Proof

Claim: There exists a Hamiltonian path in G if and only if there exists a sequence σ of length n with $Q(\sigma) = n - 1$.

Proof:

(\Rightarrow) If G has Hamiltonian path $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_n$, then sequence $\sigma = (c_1, c_2, \dots, c_n)$ has $f_{order}(\sigma) = n - 1$, so $Q(\sigma) = n - 1$.

(\Leftarrow) If sequence σ has $Q(\sigma) = n - 1$, then all n components appear exactly once and all consecutive pairs have synergy 1, meaning they correspond to edges in G , forming a Hamiltonian path.

Since DHP is NP-hard and this reduction is polynomial-time, our problem is NP-hard.

Corollary: The optimization version (finding $\max Q(\sigma)$) is NP-hard.

State Space Complexity

Search Space

The state space consists of all ordered sequences satisfying constraints:

$$S = \{\sigma \in \cup_{k=0}^n \Pi_k(C) \mid \sigma \text{ satisfies all constraints}\}$$

State space size:

$$|S| = \sum_{k=0}^n P(n, k)$$

where

$$P(n, k) = \frac{n!}{(n - k)!}$$

State Space Complexity

Theoretical Upper Bound

$$|S_{theoretical}| = \sum_{k=0}^n P(n, k) = \sum_{k=0}^n \frac{n!}{(n-k)!}$$

Asymptotic complexity: $O(n!)$

Practical

$$|S_{practical}| \approx O\left(\frac{n!}{2^d}\right)$$

03



Method Design: ToT Prompt Optimizer

Design Rationale

Why heuristic solution?

Problem: Prompt optimization is a High-Dimensional, Discrete Combinatorial Optimization problem.

Design Rationale

Why heuristic solution?

Problem: Prompt optimization is a High-Dimensional, Discrete Combinatorial Optimization problem.

Problem Challenges:

- **Combinatorial Explosion** ($O(2^N)$): Exhaustive search is intractable, requiring effective **pruning**.
- **High Evaluation Cost**: Objective function evaluation (Full Evaluation) is **costly, time-consuming, non-convex, and non-differentiable** (precluding gradient-based methods)
- **Discreteness & Constraints**: Component selection is binary/discrete, with predefined **conflict constraints**, requiring methods that support **step-wise state construction** and **validity checks**.

Design Rationale

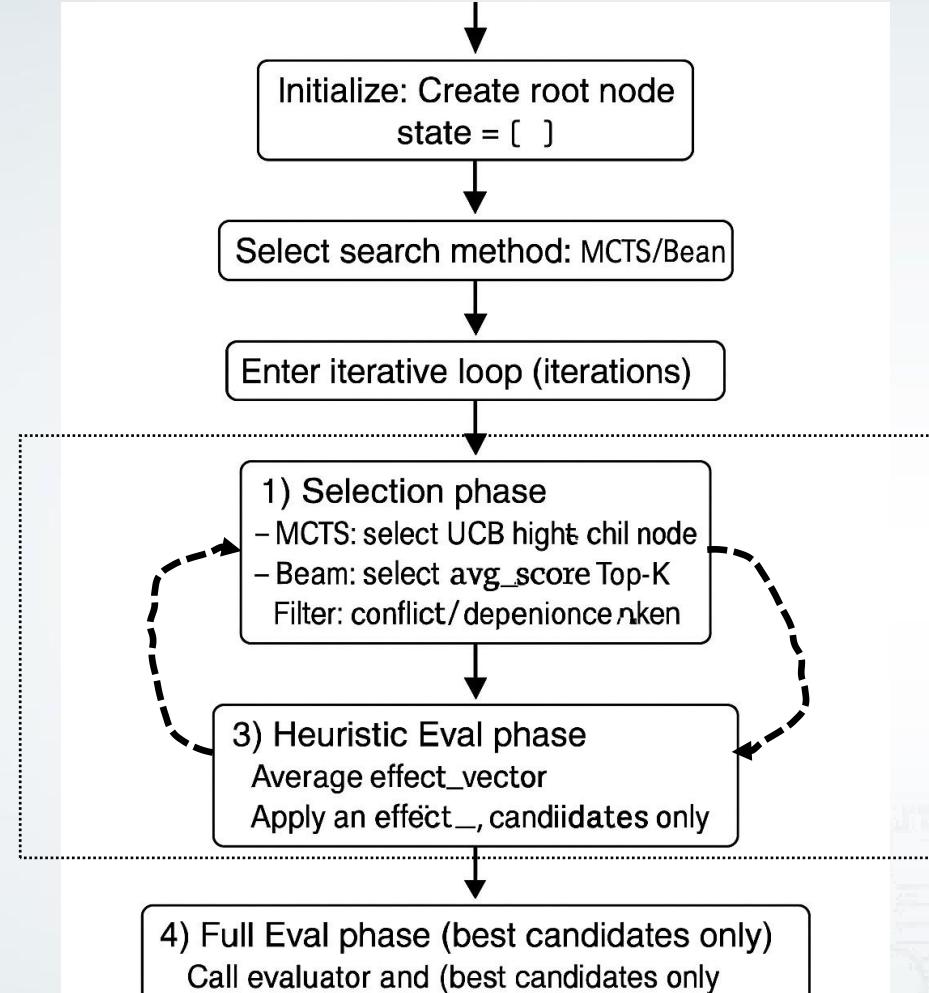
Why heuristic solution?

Justification for Metaheuristic Search

- **Monte Carlo Tree Search** (MCTS): Strategically balances exploration (new sets) and exploitation (high-performers) using **upper bound score**, ensuring costly Full Evaluations are allocated to the **most promising states**
- **Beam Search**: Offers a faster, **greedy complement** for rapid convergence to local optimums by retaining only the top k states at each step

Key Feature: Learning-to-Optimize (L2O): Incorporates a lightweight Heuristic Evaluation (predictive assessment) to guide the search.

ToT Prompt Optimizer



Core Structure and State Representation

- State: A valid set of **Prompt components**.
- Node: A position within the search tree, **storing a state S** and **search statistics**.

```
@dataclass
class SearchNode:
    """搜索树节点"""
    node_id: str
    state: List[str] # 组件ID列表
    parent_id: Optional[str]
    depth: int
    visits: int = 0
    total_score: float = 0.0
    heuristic_score: float = 0.0
    full_evaluation_score: float = 0.0
    children_ids: List[str] = None
    is_fully_evaluated: bool = False
    metadata: Dict[str, Any] = None
```

Monte Carlo Tree Search (MCTS)

MCTS guides the search through iterative sampling, executing four steps per iteration :

- **Selection:** Starting from the root, nodes are recursively selected based on **the highest heuristic score**
- **Expansion:** Valid candidate child states (enforcing **conflict avoidance** via `generate_candidates`) are generated from the selected node. New child nodes are created and immediately undergo Heuristic Evaluation (`evaluate_state_heuristic`).
- **Evaluation** (Simulation/Rollout): The most promising nodes (e.g., `avg_score > best_score × 0.7`) are subjected to a costly Full Evaluation (**`evaluate_state_full`**) to obtain the accurate `full_score`.
- **Backpropagation**

Heuristic Scoring Function

Purpose: Provide a **lightweight approximation** of prompt quality to guide the search efficiently, avoiding expensive full LLM evaluations at every node.

Formula: Weighted sum of five metrics $\in [0, 1]$:

$$H = \sum_{m \in M} w_m \cdot S_m$$

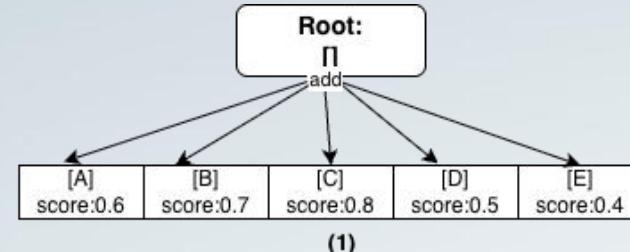
- $M = \{\text{Faithfulness}, \text{Conciseness}, \text{Completeness}, \text{Readability}, \text{Insightfulness}\}$
- Pre-defined Weights

Metric Calculation Strategy: Balances speed and reliability

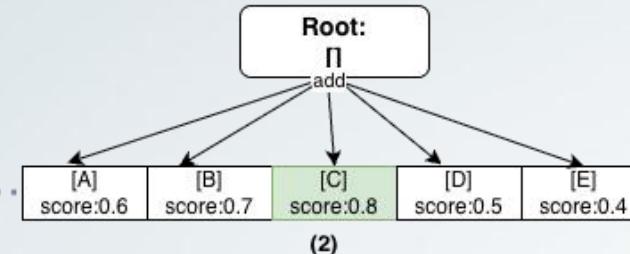
Rule-Based Metrics (Fast Heuristics): **Conciseness** on word count, **Completeness** on keyword coverage, **Readability** on average sentence length

LLM-Based Metrics (Accurate Heuristics): Faithfulness and Insightfulness

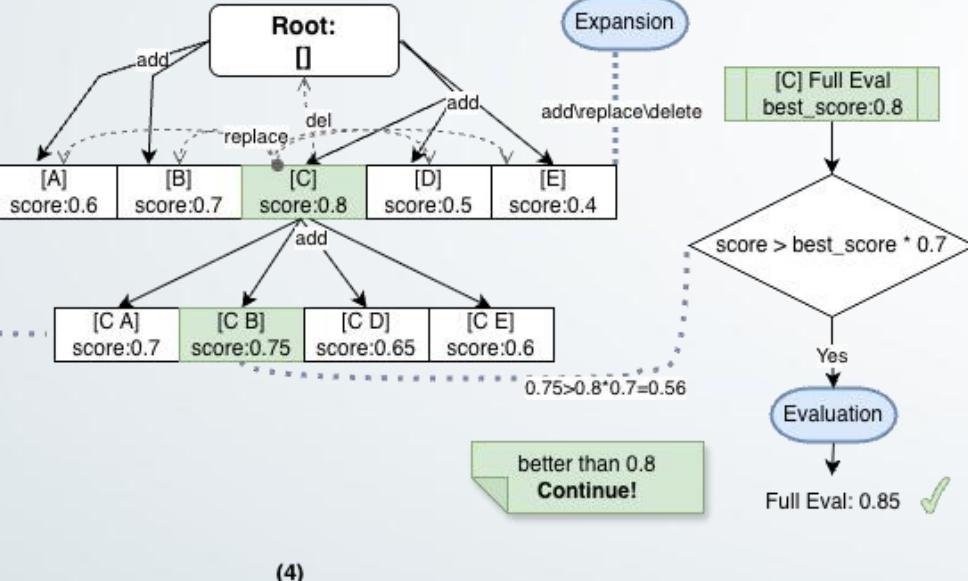
Given 5 non-conflict components: [A B C D E]



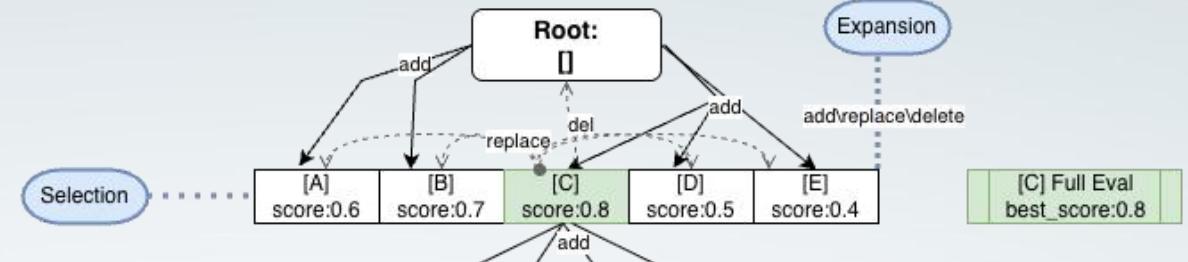
Π Full Eval
best_score:0



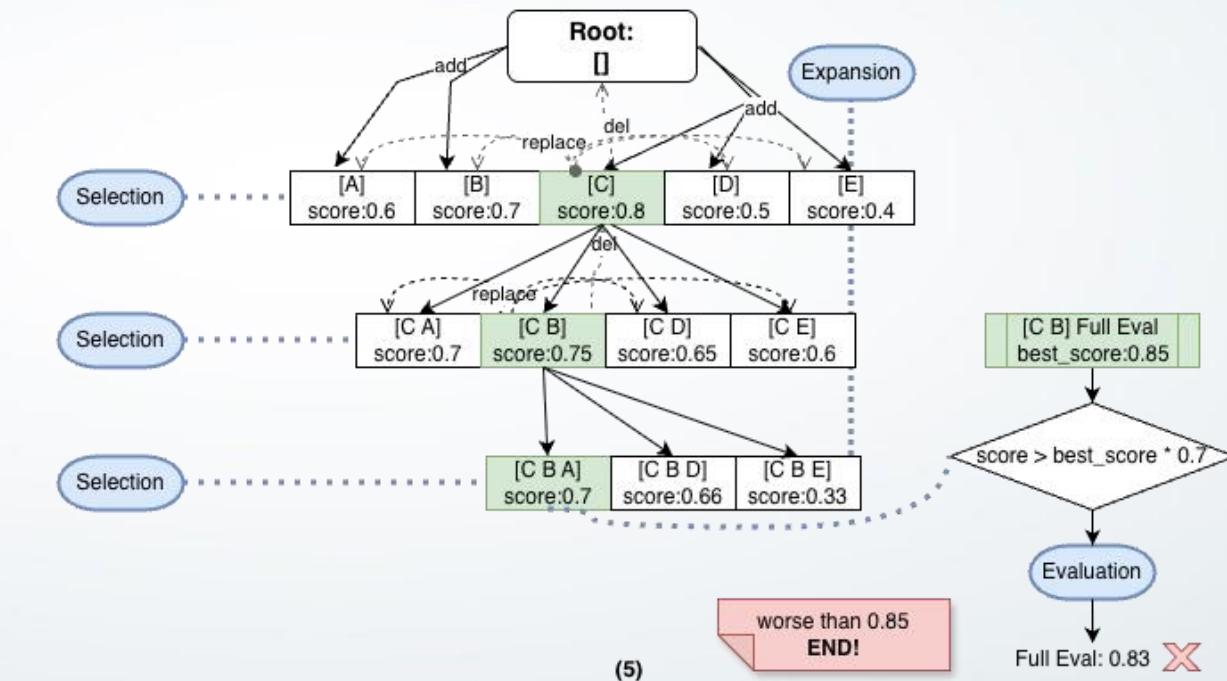
[C] Full Eval
best_score:0.8



score > best_score * 0.7
Yes
Evaluation
Full Eval: 0.85 ✓



[C] Full Eval
best_score:0.8



Full Eval: 0.83 X

Beam Search

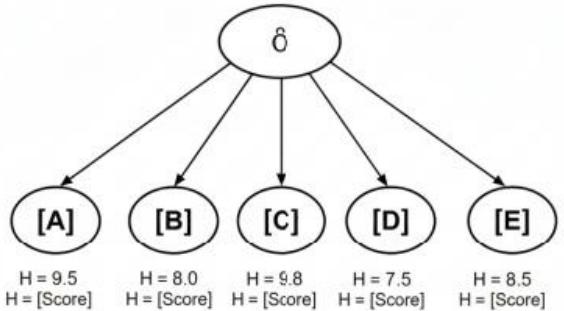
Beam Search implements a more **greedy** heuristic approach :

- **Breadth-First Pruning:** In each iteration, the **top-k** scoring nodes (the beam) from all existing nodes are selected.
- **Expansion and Heuristics:** All **new candidate states** are generated from the k beam nodes and assessed using Heuristic Evaluation.
- **Evaluation:** The new candidates are ranked by their Heuristic Score, and **only the top k** are selected to receive a Full Evaluation. This ensures a fast, concentrated search.

Given 5 non-conflict components: [A B C D E]

Step 1: Depth 1 Expansion & Selection

Depth 0



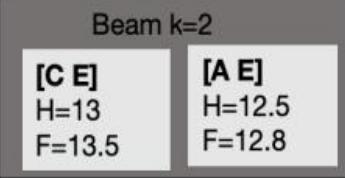
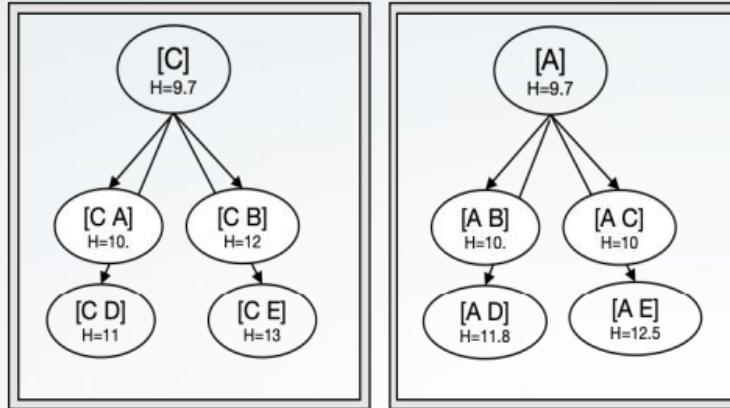
Current Beam k=2



Legend
H = Heuristic Train Score
F = Full Eval score

(1)

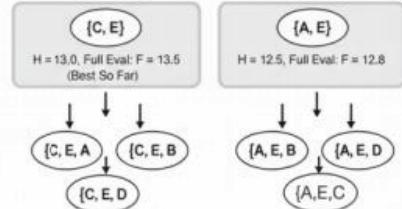
Step 2: Depth 2 Expansion & Selection



(2)

Step 3: Depth 3 Expansion & Selection (Beam k=2)

Depth 2 Current Beam



6 Total New Candidates



Final Beam k=2



(3)

Per-State Operations

- Constraint Checking (Dependency, Mutual Exclusion, Length, Position): Total complexity $O(k)$ with optimization, where k is the sequence length
- Quality Evaluation $O(\sigma)$: $O(k)$ for component base score, order synergy, and penalty terms

Algorithm	Per Iteration/Layer Complexity (General)	Worst Case Asymptotic Complexity
MCTS	$O(I(d \cdot \log b + b + C_e))$	$\mathcal{O}(I \cdot n)$
Beam Search	$O(d \cdot w \cdot b \cdot \log(w \cdot b))$	$\mathcal{O}(n^2 \log n)$

Search Algorithm Complexities

- Variables: I = iterations, d = depth, b= branching factor (avg.), C_e = evaluation cost, w = beam width, n = total components
- Efficiency: Both MCTS and Beam Search are significantly more efficient than the theoretical $O(n!)$ brute-force search space

Algorithm	Per Iteration/Layer Complexity (General)	Worst Case Asymptotic Complexity
MCTS	$O(I(d \cdot \log b + b + C_e))$	$O(I \cdot n)$
Beam Search	$O(d \cdot w \cdot b \cdot \log(w \cdot b))$	$O(n^2 \log n)$

04

Experimental Results and Discussion

Experiments

Dataset → Summarization Task

Instead of complex formats (URLs, PDFs) or dialogue-style content.

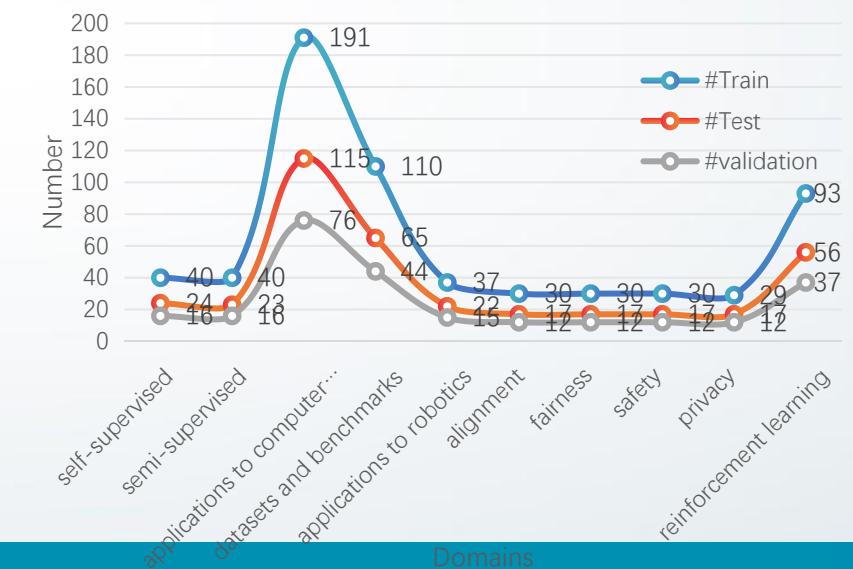
A clean, plain-text dataset was manually constructed for the experiments.

Source:

- Summarization: ICLR 2026 OpenReview platform
- Prompt Components: Human designed 20 Entries

Table 5: Statistics of word length across dataset splits.

Category	Content Length	Summary Length
Training Dataset	182.82	26.94
Test Dataset	184.56	26.50
Validation Dataset	174.90	26.35



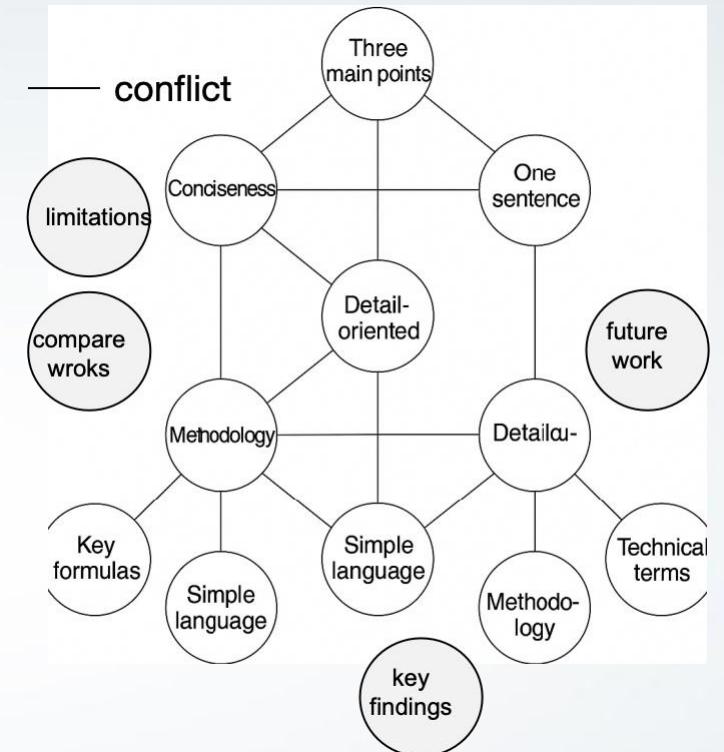
```
{
  "paper_id": "bU8tRjuanU",
  "title": "Low-Rank Attention and Contrastive Alignment for Deep Multi-View Clustering",
  "domain": "self-supervised",
  "content": "Recent years have witnessed significant advancements in deep multi-view clustering (MVC). However, prevailing methods exhibit three critical limitations: (1) poor scalability for large-scale datasets, (2) neglect of anchor semantic consistency in feature alignment, and (3) inability to capture high-order feature interactions. To overcome these challenges, we propose a Low-Rank Attention and Contrastive Alignment framework (LRACA). Unlike conventional approaches that align sample-level features in shared subspaces, LRACA employs a category-aware anchor generation module to directly align high-level semantic prototypes (i.e., category centers) across views, explicitly enforcing clustering semantic consistency. Furthermore, we devise a dynamic low-rank attention mechanism to enhance feature discriminability, where entropy regularization constrains attention weight distributions to derive clustering pseudo-labels. Finally, a pseudo-label-guided cluster-level contrastive learning module maximizes cross-view mutual information through a feed-forward optimization paradigm. Extensive experiments on six large-scale multi-view datasets demonstrate that LRACA significantly outperforms state-of-the-art methods.",
  "gold_summary": "This paper proposes a deep multi-view clustering algorithm equipped with an anchor-based attention mechanism and cluster-based contrastive learning, which achieves good performance, but has many problems, as detailed in the weaknesses section."
}
```

(a) An Example of Data Entry

```
{
  "id": "conciseness",
  "text": "Please use concise language",
  "description": "Requires concise output",
  "effect_vector": {
    "conciseness": 0.8,
    "readability": 0.6,
    "completeness": -0.2
  },
  "conflicts": [
    "detail_oriented",
    "comprehensive"
  ],
  "requires": [],
  "estimated_tokens": 8
},
```

(b) An Example of Prompt Component Entry

Fig. 4: Data Entry Statistic.



Setup

Models:

- Base Model (for ToT Training): Qwen3-1.7B.
- Scoring Model (Discriminator): Qwen3-8B

Baselines:

- **SimplePrompt**: “Summary this paper”
- **DetailedPrompt**: “A prompt asking for core content, methodology, and contributions”
- **StructuredPrompt**: “A prompt asking to summarize from four aspects: background, methodology, experiments, and conclusions”

Metrics:

- **N-gram Static Metrics**: ROUGE-2 F1 and ROUGE-L F1
- **Semantic Similarity Metrics**: BERTScore F1 (using "bert-base-multilingual-cased").
- **LLM Evaluation Score**: Heuristic Scoring

Experimental Results and Discussion

Overall Performance

- **ToT Prompt-Beam's Success:** This strategy consistently achieved the highest scores across all reported metrics
- MCTS vs. Beam Search:

ToTPrompt-MCTS < **ToTPrompt-Beam**

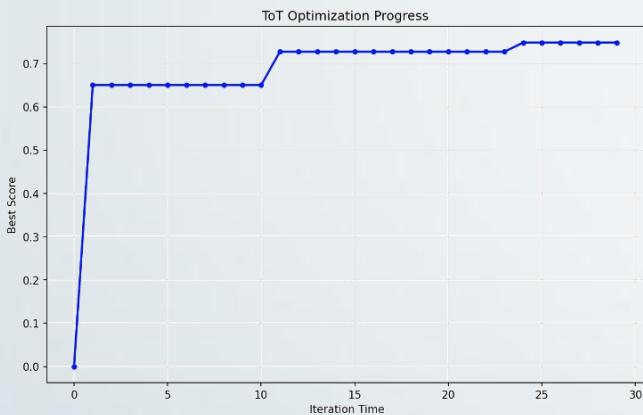
Prompt Strategy	Heuristic Score ↑	ROUGE-2 ↑	ROUGE-L ↑	BERTScore ↑
SimplePrompt	0.51 ± 0.10	0.03 ± 0.02	0.03 ± 0.02	0.05 ± 0.13
DetailedPrompt	0.58 ± 0.06	0.02 ± 0.01	0.07 ± 0.02	-0.01 ± 0.04
StructuredPrompt	0.58 ± 0.04	0.01 ± 0.01	0.06 ± 0.01	-0.18 ± 0.20
ToT-MCTS(ours)	$0.61 + 0.00$	0.03 ± 0.02	0.09 ± 0.01	$0.06 + 0.04$
ToT-Beam(ours)	0.67 ± 0.02	0.07 ± 0.05	0.13 ± 0.05	0.16 ± 0.06

Convergence Analysis

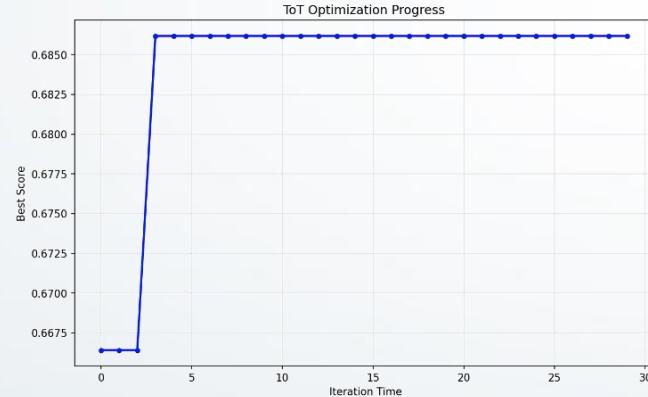
Scores quickly converged in both strategies. Beam Search converged significantly faster than MCTS due to its greedy-like selection at each step

Potential Reason for Rapid Convergence

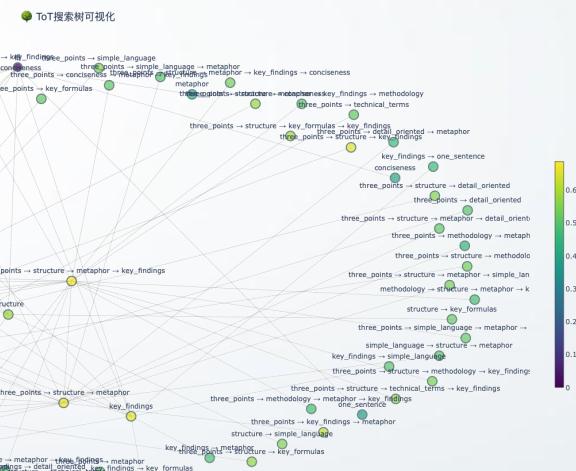
- an insufficiently rich set of prompt components,
- a small candidate set
- a small training batch size leading to minimal score variation



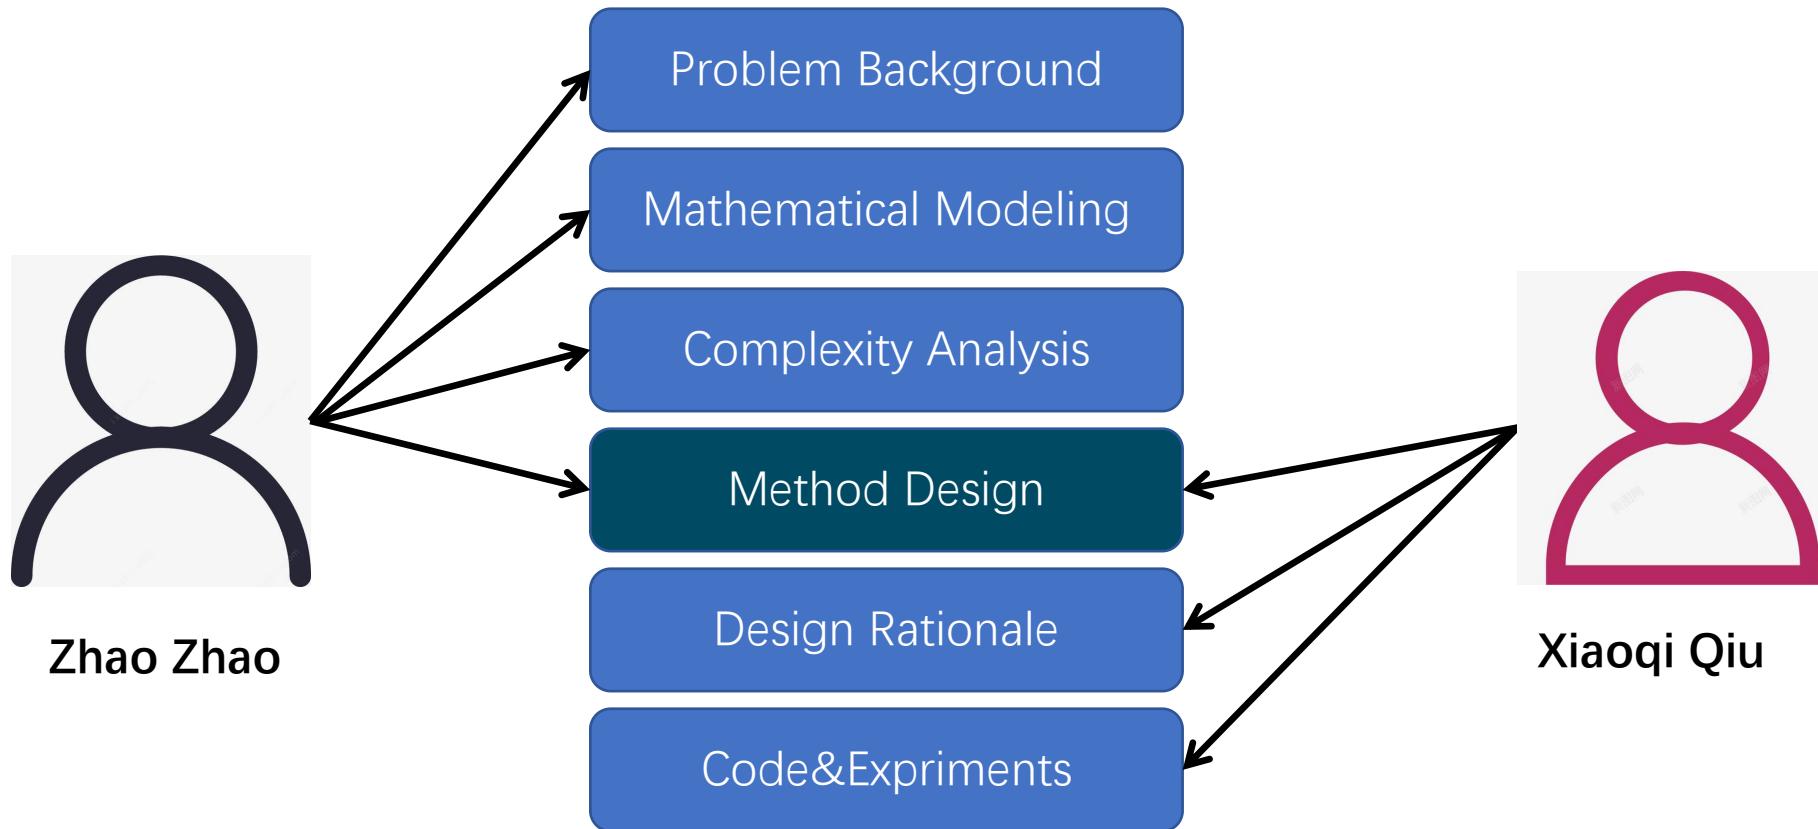
(a) Optimization Progress of MCTS Search



(b) Optimization Progress of Beam Search



Division of Labor



Code&Dataset:
git@github.com:Siki-cloud/ToT_Optimizer_for_LLMSum.git

“

ToT-Prompt Optimization for LLM-based Automatic Summarization

Thank you

”

Reporter: Zhaozhao Xiaoqi Qiu

Date: 2025/12/10

